

Foundations of Formal Market Microstructure in U.S. Equities

Machine-Checked Allocation, Composition,
Necessary State, and Regulatory Traceability

Miquel Noguer i Alonso
Artificial Intelligence Finance Institute

September 5, 2026

DOI: [10.5281/zenodo.22343804](https://doi.org/10.5281/zenodo.22343804)

GitHub: [MiquelNoguerAlonso/Formal-Verification](https://github.com/MiquelNoguerAlonso/Formal-Verification)

Abstract

This paper develops executable models of price–time priority and parity allocation in U.S. equities to establish when execution consistency holds and which state must be retained. Splitting an order can change who receives shares. In a parity wheel, resetting the rotation between executions can alter allocations even when total incoming quantity is unchanged.

A counterexample shows that restarting the wheel between executions can change allocations. Preserving its participant pointer and unfinished lot allowance restores consistency between split and combined quantities on reachable states, under a positive lot size and an admissible first quantity. A further invariant recovers the unused allowance from the cumulative allocation of the current participant. Consequently, when the pointer and aligned cumulative fills are known, the allowance need not be stored independently.

The information results distinguish prediction from reconstruction. On specified nondegenerate families sharing an authenticated observation, every summary sufficient to predict future fills must distinguish live pointer positions. The pointer is sufficient once the family is fixed, and a reachable family attains the exact state bound. Separate finite families characterize the additional information required when cumulative fills are unavailable.

Supporting results establish allocation feasibility, conservation, priority, one-lot balance, and explicit approximation bounds relating wheel allocations to constrained equal awards. They also yield a conventional analytic limit as lot size vanishes. A decidable price–time checker is characterized by feasibility and seriality, with conformance preserved under sequential execution.

Lean 4.22.0 checks 362 theorem declarations using its core library. Axiom dependencies are audited explicitly; the analytic limit is distinguished from the formalized results. A dated regulatory ledger links 58 selected provisions or identified source items to executable definitions and audit records. These mappings support inspection without proving legal interpretation, evidence authenticity, or equivalence to a production exchange. The resulting framework connects allocation mechanisms, persistent state, and reproducible conformance analysis.

Keywords: Market microstructure; U.S. equities; formal verification; Lean 4; price–time priority; parity allocation; stream consistency; state identifiability.

Contents

1	Introduction	4
1.1	The claim	4
1.2	What is built	4
1.3	Proof system	5
1.4	How to read this paper without advanced prerequisites	5
1.5	Related work	6
2	Primitive objects	6
3	The allocation operators	7
4	The algebra of operators	8
5	Axioms for allocation	9
6	Characterization theorems	9
7	Structure theorems	11
8	Allocation as rationing, and the pointer theorem	12
8.1	The identification	12
8.2	What composes, and on which index	12
8.3	The pass-reset wheel does not	13
8.4	The pointer-and-lot wheel does	13
8.5	What the theorem establishes	14
9	Cross-market structure	15
9.1	Consolidation	15
9.2	Order protection	15
9.3	Tick, lock, fee	16
9.4	Two market-data clocks	16
9.5	Order types	16
9.6	Bands and tests	17
10	Auction structure	17
11	Observability structure	18
12	Conformance	21
13	Rule-by-rule coverage register	21
14	Regulatory traceability: audit contract	24
15	Scope and limitations	25
16	Conclusion	25

A	The proof system	26
A.1	Load-bearing theorem contracts	26
A.2	Paper and source availability	28
B	Dated regulatory traceability and inspection	29
B.1	Fixed and dated corpus	29
B.2	Formalization of non-numerical duties	29
B.3	Federal specifications	30
B.4	Plan and venue specifications	31
B.5	Ledger audit and exact force	32
B.6	Inspection contract	32
C	Compiler-checkable source excerpt	33
D	Certified inequalities for the strategic layer	34
D.1	Contests for a slot or a tranche	34
D.2	The race for setter priority	34
D.3	Privileged slot and venue competition	34
	References	35

1 Introduction

1.1 The claim

An exchange receives marketable quantity as a stream. If 200 shares arrive at once or as two consecutive orders of 100, a stream-consistent allocation mechanism should reach the same final state after processing the same total quantity. Price-time has this property using residual claims alone. A parity wheel does not: if every execution restarts at the first participant, order segmentation changes who trades.

The paper’s central contribution is a state-and-information result. First, an executable counterexample proves that reset rotation violates Moulin composition. Second, a wheel that persists its participant pointer and unfinished round-lot phase satisfies exact split-versus-combined consistency on every reachable state and every admissible first estate. Third, for any indexed nondegenerate family with the same authenticated observation, future sufficiency forces a summary to encode the pointer injectively, and the pointer is itself a sufficient code once the family is fixed. A reachable full-allowance family realizes the exact m -state bound. Finally, on every reachable live-head state, the unused allowance is $L - (A_p \bmod L)$. Consequently, when participant-aligned cumulative fills are available, the allowance is derived rather than separately stored and the pointer label is the remaining phase datum in the stated observation model.

Market microstructure theory studies the economic consequences of trading rules (O’Hara, 1995; Parlour, 1998; Harris, 2003). Here executable allocation operators are the primitive objects. The paper also isolates the fixed-order characterization of price-time, proves one-lot balance and approximation to constrained equal awards, derives a vanishing-lot continuum limit, and turns the price-time characterization into a deterministic conformance check. These results concern the specified mechanisms. Correspondence with institutional rule text remains a modelling claim, and real-world compliance remains an evidentiary claim.

The selected mechanics use integer shares, round lots, tick-constrained prices, rotation, minimum, comparison, and lexicographic maximization. Their finite invariants are proved by structural induction and exact arithmetic in Lean. The continuum-limit corollary is a conventional real-analysis deduction from the machine-checked finite error bound. The strategic propositions in Appendix D concern best responses in contests, rent dissipation, dilution of a privileged slot, and routing by net cost under additional behavioral assumptions.

1.2 What is built

Section 2 fixes the primitive objects. Section 3 defines the three allocation operators. Section 4 gives their algebra: sequential composition, conservation closure, benchmark naturality, relabeling equivariance of the persistent wheel, and monotonicity. Section 5 states the allocation predicates and Section 6 proves uniqueness (seriality is price-time), a one-lot balance guarantee for the wheel, incompatibility, and ladder uniqueness. Section 7 adds completeness, conservation, displayed precedence, and price priority across levels. Section 8 separates estate composition, water-level composition, and state-transition composition, proves feasibility of the estate-indexed unit rule, isolates the reset-composition obstruction, and proves that the persistent pointer wheel is both balanced and stream-consistent under stated hypotheses. Sections 9–10 give cross-market and auction kernels. Section 11 proves non-identifiability, general fiberwise pointer minimality, and exact predictive-summary bounds on reachable finite families, with separate contracts for observers who already know cumulative fills. Section 12 turns price-time uniqueness into a complete, decidable and composition-safe check. Section 13 registers every analytic rule family. Section 14 gives the regulatory audit contract; Appendix B contains the dated clause ledger, inspection profiles, and

source pinpoints. Section 15 separates theorem, traceability, and production claims. Appendix D isolates optional strategic inequalities and their additional behavioral assumptions.

1.3 Proof system

The executable mechanisms and named theorem contracts are formalized in Lean 4 (de Moura and Ullrich, 2021) version 4.22.0 with the core library only—no Mathlib (The mathlib Community, 2020)—as sixteen proof/test modules (including the axiom-audit module) plus one command-line checker entry point. `Conformance.lean` contains deterministic test harnesses; `AllocationRecord.lean` is the checker entry point. Each formal result below is annotated with its Lean declaration; fixed-length bit-encoding consequences in Section 11 are identified explicitly as mathematical corollaries of checked injectivity. The package builds with `lake build`; `Check.lean` prints the axiom dependency of all 362 theorem declarations; `lake exe allocation_record` runs the checker on a supplied order-level queue and queue-aligned fill vector. The source has no `noncomputable` declaration and no use of `Classical.choose`; executable definitions therefore remain distinct from the 48 theorem proof terms whose audited dependencies include `Classical.choice`. Appendix A names those declarations and gives exact contracts for load-bearing results.

1.4 How to read this paper without advanced prerequisites

No programming or proof-theory background is required to understand the economic claims. The formal notation makes each rule precise, but the central objects have direct market meanings.

The market picture. A *resting order* is an instruction already waiting at a venue. A new, marketable order arrives with a quantity to buy or sell. An *allocation rule* decides how much of that incoming quantity each resting order receives. The resulting executed amounts are *fills*. A *queue-aligned allocation record* associates those fills with the supplied resting-order positions; unlike a public trade report, it is privileged audit data. A *queue* is the ordered list of resting order sizes at one price.

The notation. $\mathbb{N}$ means the whole numbers $0, 1, 2, \dots$. A list such as $Q = (200, 300, 500)$ describes three resting orders. If $x = 600$ shares arrive, an allocation such as $A = (200, 300, 100)$ says that the first order receives 200 shares, the second 300, and the third 100. The symbol $\min\{a, b\}$ means “whichever number is smaller,” $\sum$ means “add the entries,” and $\lfloor z \rfloor$ means “round down to a whole number.” A function is just a rule that turns specified inputs into an output.

The formal-verification picture. Lean is software that checks a proof one logical step at a time. A Lean theorem is not a simulation and not a statistical confidence statement: if its assumptions are satisfied, its conclusion follows exactly. The clause layer prevents a named provision from disappearing silently, but it cannot decide whether a venue supplied truthful evidence for a procedural field. Clause coverage, theorem validity, and real-world compliance are therefore three different claims.

How to read a theorem. Read each result as “under these conditions, this guarantee follows.” Text in typewriter font, such as `priceTime_sum_le`, is the name of the matching Lean declaration; it can be ignored on a first reading. Proof ideas explain why a result is true, while the full machine-checked proof is in the accompanying source. A counterexample proves that a broad claim cannot hold in general by exhibiting one valid case in which it fails.

Abbreviations. NBBO is the national best bid and offer; NBB and NBO are its bid and offer components. SIP is a securities information processor; ISO is an intermarket sweep order; DMM is a designated market maker; SRO is a self-regulatory organization; ATS is an alternative trading system; LULD is the Limit Up–Limit Down Plan; NMS is the national market system; and Reg SHO is Regulation SHO.

A running example. Suppose the queue is (200, 300, 500) and 600 shares arrive. Price–time gives (200, 300, 100) because it finishes earlier orders before later ones. A parity wheel with a 100-share round lot can give (200, 200, 200) by visiting each participant repeatedly. Both allocate 600 shares without exceeding any resting order, but they express different ideas of priority and balance. Much of the paper formalizes that difference.

1.5 Related work

Formal work on exchange mechanisms is directly relevant. [Sarswat and Singh \(2020\)](#) formalize fairness, uniformity, and individual rationality for financial-market trades; [Natarajan et al. \(2021\)](#) extend verified double auctions to multiple quantities, prove uniqueness results, extract programs, and test them on market data; and [Garg and Sarswat \(2022\)](#) specify continuous double auctions, prove that three properties determine their input–output relation, and extract a verified trade-log checker. Verified financial software also includes the Imandra dark-pool analysis of [Passmore and Ignatovich \(2017\)](#) and its later exchange digital-twin work ([Brennan, 2024](#)), as well as contract languages ([Bahr et al., 2015](#)); formal law includes Catala ([Merigoux et al., 2021](#)). The present contribution is narrower and complementary: machine-checked integer models of price–time allocation, NYSE-inspired setter/parity allocation, their rationing identities, and selected cross-market rule fragments. It does not claim the first formalization of exchange matching or automated conformance checking.

The round-robin structure of parity is related to fair queueing ([Nagle, 1987](#)). The optional strategic layer connects to queue rationing under price controls ([Yao and Ye, 2018](#)), fee structure ([Comerton-Forde et al., 2019](#)), fragmentation ([O’Hara and Ye, 2011](#)), latency ([Ding et al., 2014](#); [Bartlett and McCrary, 2019](#)), frequent batch auctions ([Budish et al., 2015](#)), and contest models ([Tullock, 1980](#); [Baye et al., 1996](#)). The axiomatic theory of rationing is due to [Young \(1987\)](#), [Young \(1994\)](#), and [Moulin \(2000\)](#), and is surveyed by [Thomson \(2003\)](#); [Aumann and Maschler \(1985\)](#) provide the contested-garment lineage, [Herrero and Villar \(2001\)](#) study classical divisible rules, and [Herrero and Martínez \(2008\)](#) study balanced claims allocation with indivisible units. The composition axiom is used here in the formulation emphasized by [Moulin \(2000\)](#).

2 Primitive objects

Definition 2.1 (Quantities). A share count is an element of $\mathbb{N}$. A round lot is a positive integer L . A price is an element of $\mathbb{N}$ in units of \$0.0001; we write `dollar` = 10^4 , `penny` = 100.

Definition 2.2 (Queue). A queue at a price is a finite participant-labelled list in a canonical external order:

$$Q = ((i_1, q_1), \dots, (i_m, q_m)).$$

Serial allocation uses arrival order. The pointer-wheel state stores a rotated internal view whose head is the current pointer; statements about participant allocations compare outputs only after mapping that rotated view back to the canonical labels.

Definition 2.3 (Level, ladder). A level is a price with a queue; on a hybrid venue it also carries a setter tranche Q^s (**Level**). A ladder is a list of levels ordered by price.

Definition 2.4 (Participant state). A participant is $(a, r) \in \mathbb{N}^2$: allocated so far, remaining. A level in state form is a list of participants; $a + r$ is invariant under allocation and equals the participant's original size.

Definition 2.5 (Quote). A quote is (venue, price, size, automated) (**Quote**). Within the model, it is protected iff automated and at least one round lot (**protected_**). Actual Regulation NMS protection also requires the quotation to be the exchange's best bid or offer; the model assumes its input has already been restricted to venue-top quotations ([U.S. Securities and Exchange Commission, 2005](#)).

Definition 2.6 (Fill, allocation record, market-data event). A fill is an executed size. A queue-aligned allocation record is the vector of fills matched to the supplied order positions; it is an audit input and is not asserted to be a public tape. Separately, a market-data event carries a venue timestamp v and a latency $\ell \geq 0$ and is consolidated at $v + \ell$ (**Ev**). The two-clock results concern only these timestamps, not participant-level attribution.

Definition 2.7 (Allocation, operator). An allocation of x over Q is a list $A \in \mathbb{N}^m$. An operator is a function $x \mapsto (A, x')$ giving fills and leftover (**Op**).

For equal-length allocation lists, write $A \oplus B := \text{List.zipWith}(\text{funab} \Rightarrow \mathbf{a} + \mathbf{b}) \mathbf{AB}$ for componentwise addition. Every displayed sum of allocation vectors below uses $\oplus$.

Pointwise relations between lists of equal length are written with the inductive predicate $\text{Pw } P$ ($\text{Pw.nil}, \text{Pw.cons}$).

Plain-language reading. The symbols summarize one transaction. Q is what was available, x is what arrived, A is what traded against each resting order, and x' is any quantity still unfilled. A correct rule must never fill more than was available and must account for every share.

3 The allocation operators

Definition 3.1 (Serial: price-time; **priceTime**). $a_i = \min\{q_i, x_{i-1}\}$ with $x_0 = x$, $x_i = x_{i-1} - a_i$.

Definition 3.2 (Parity: visit, pass, wheel; **give**, **wheelRound**, **wheel**). A visit gives **give** $L r x = \min\{L, \min\{r, x\}\}$. A pass visits $r_1, \dots, r_m$ in order, each visit reducing the quantity. The wheel iterates passes, subtracting fills from sizes, until quantity is exhausted or a fuel bound n is reached; it returns cumulative fills and leftover. The state-form wheel **wheelS** does the same on participants (a, r) via **visit** and **passS**; **initS** builds the initial state from sizes.

Definition 3.3 (Hybrid: setter-then-parity; **rule737**, **rule737full**). Tier 1 gives $t_1 = \min\{Q^s, x\}$ to the setter. Tier 2 runs the wheel on displayed sizes with $x - t_1$. Tier 3 runs the wheel on non-displayed sizes with the tier-2 leftover.

Definition 3.4 (Integer equal-awards benchmark; **waterFill**). At water level s , $W_i = \min\{q_i, s\}$.

Here $s \in \mathbb{N}$: this is a discrete constrained-equal-awards benchmark, not a real-valued continuous limit. Its relation to the wheel is Theorem [6.4](#).

Example 3.5 (A hybrid allocation; `ex_rule737`, `ex_waterfill`, `ex_within_L`, `ex_book_share`, `ex_priceTime`). $L = 100$. Four participants—a market maker, two brokers, and an aggregated electronic book—display (300, 500, 200, 1000); the second broker (third in the list, displaying 200 shares) holds setter priority on that 200; the pointer is at the market maker; $x = 1300$. The kernel evaluates `rule737 100 20 200 (300, 500, 0, 1000) 1300 = (200, (300, 400, 0, 400), 0)`: setter 200, then wheel fills 300, 400, 0, 400, nothing left. Constrained equal awards at $s = 400$ on the post-setter parity pool gives (300, 400, 0, 400) and clears 1100: wheel and benchmark coincide here. Within that parity pool, the electronic book holds $1000/1800 \approx 55.6\%$ of depth and receives $400/1100 \approx 36.4\%$ of parity fills ($400 \cdot 1800 < 1000 \cdot 1100$). On a serial venue with arrival order (200, 300, 500, 1000) the fills are (200, 300, 500, 300). All five statements are proved by `decide` and depend on no axioms.

Plain-language reading. The three mechanisms answer the same practical question in different ways. Price–time rewards earlier arrival. Parity repeatedly shares round lots among eligible participants. The hybrid first reserves a specified setter tranche and then applies parity to the remaining pools. The numerical example lets a reader compare their consequences without reading code.

4 The algebra of operators

Definition 4.1 (Sequential composition; `seq`, `idOp`, `ladderOp`). $(f \triangleright g)(x) = (A_f \mathbin{++} A_g, x'_g)$ where $(A_f, x'_f) = f(x)$ and $(A_g, x'_g) = g(x'_f)$. The identity allocates nothing. A ladder is the right fold of its level operators under $\triangleright$.

Definition 4.2 (Conservation; `Conserves`). f is conservative if $\sum A + x' = x$ for every x .

Theorem 4.3 (Monoid and conservation closure; `seq_assoc`, `seq_id_left`, `seq_id_right`, `seq_conserve`, `ladderOp_conserve`). *Operators form a monoid under $\triangleright$ with identity `idOp`. Conservative operators are closed under composition, so every finite ladder assembled from conservative stages is conservative.*

Theorem 4.4 (The hybrid is a composite; `setterOp`, `wheelOp`, `rule737Op`, `rule737Op_eq`, `rule737Op_conserve`). *After flattening the three tier allocations into one fill list, `rule737full` agrees with `setter` $\triangleright$ `wheelR` $\triangleright$ `wheelH`, including its leftover. Its conservation law follows from Theorem 4.3: `setterOp_conserve` certifies the setter, and `wheelOp_conserve` is applied once to the displayed parity tier R and once to the non-displayed parity tier H . The separately audited `wheelOpFull_conserve` certifies the canonical-fuel wrapper; `rule737Op` itself is the explicit-fuel composite.*

The theorem is architectural. A rule change that inserts a tier, reorders tiers or adds a level changes a composite, and the invariants closed under composition need not be re-proved for the new rule.

Theorem 4.5 (Equivariance under list transformations). Lean: `waterFill_map` and `priceTime_not_anonymous`.

Let $\sigma : \text{ListNat} \rightarrow \text{ListNat}$ satisfy the `List.map` naturality equation

$$\sigma(\text{List.map } f \ell) = \text{List.map } f (\sigma \ell)$$

for every $f : \mathbb{N} \rightarrow \mathbb{N}$ and every list ℓ . Then `waterFill s` (σQ) = σ (`waterFill s` Q). List permutations satisfy this hypothesis. Price–time is order-sensitive: on the queue (200, 300) with $x = 200$, reversing the queue and reversing the fills do not commute.

Theorem 4.6 (Relabeling equivariance of the persistent wheel; `runW_relabel`, `runWFull_relabel`). *Let f relabel participant identifiers while preserving every allocation, remaining claim, cyclic position, and current allowance. Relabeling before a canonical wheel execution gives exactly the relabeling of its returned state, with the same leftover. For a bijective f , this is equivariance under renaming participants with the pointer transported in the same rotation. It does not permit rearranging priority positions independently of the state.*

Theorem 4.7 (Monotonicity; `priceTime_mono`, `waterFill_mono`). *Price-time is pointwise monotone in incoming quantity; water-filling is pointwise monotone in the water level.*

Permutation equivariance is one anonymity benchmark; it is neither a characterization of venue parity nor, by its failure alone, a characterization of time priority.

5 Axioms for allocation

Axiom 5.1 (Feasibility; `Feasible`). $a_i \leq q_i$ for all i and $\sum_i a_i = \min\{x, \sum_i q_i\}$.

Axiom 5.2 (Seriality; `Serial`). For every i , if $\sum_{j>i} a_j > 0$ then $a_i = q_i$.

Axiom 5.3 (Band; `Band`). There is a water level w with $a_i \leq w + L$ for all i and $a_j \geq w$ for every unexhausted j .

These three are allocation predicates that an operator has or lacks; the next section says which. Displayed-before-hidden is instead a proved structural property of the particular three-tier operator and appears in Section 7.

6 Characterization theorems

Remark 6.1 (One family, two endpoints; `wheelRound_eq_priceTime_of_large_lot`). If $L \geq \max_i q_i$, then one pass of the round-lot wheel equals `priceTime` on every estate x . Price-time is therefore the large-lot endpoint of the pass family. Repeating the same pass primitive with smaller L gives the finite-increment wheel studied below; a single small-lot pass need not allocate the entire feasible estate.

Theorem 6.2 (Uniqueness; `priceTime_feasible`, `priceTime_serial'`, `serial_unique`). *Price-time satisfies Axioms 5.1 and 5.2, and every allocation satisfying both equals price-time.*

Proof idea. Induction on the queue. If anything behind the head is filled, the head is full and $x \geq q_1$; otherwise the head carries $\min\{x, \sum q\}$, which is x when $x < q_1$. Either way $a_1 = \min\{q_1, x\}$, and the tail satisfies both axioms with $x - a_1$. $\square$

This is an elementary characterization for a fixed priority order. [Moulin \(2000\)](#) already treats indivisible units and characterizes the family of priority rules through consistency, upper composition, and lower composition. The present hypotheses instead directly impose seriality and feasibility on one allocation. The contribution here is the executable list operator, its machine-checked contract, and the conformance corollary; indivisibility itself is not a new extension of Moulin's setting.

Theorem 6.3 (Balance; `AllAt`, `passS_band`, `passS_eq`, `wheelS_band`, `wheelS_balanced`). *For all L , fuel n , sizes and x , the wheel's final state satisfies Axiom 5.3. Consequently no participant is more than L ahead of any unexhausted participant (*Balanced*).*

Proof idea. Invariant **AllAt** v : every allocation is $\leq v$ and every unexhausted allocation is $= v$. Any pass maps **AllAt** v into **Band** v , because a visit adds at most L and an unexhausted participant was unexhausted before. An un-truncated pass maps **AllAt** v to **AllAt** $(v + L)$, because while quantity remains every visit gives exactly $\min\{L, r_i\}$ (**give_of_pos**, **wheelRound_full**). Induction on fuel. $\square$

Theorem 6.4 (Within one round lot of constrained equal awards; **wheelLevel**, **wheelS_within_L_explicit**, **wheel_within_L**). *For each run, the executable value $w = \text{wheelLevel } L \ n(\text{initS } Q) \ x \ 0$ adds L after each pass that leaves positive incoming quantity, and every participant satisfies $|a_i - \min\{Q_i, w\}| \leq L$. The statement holds for the state-form wheel and, by **wheelS_eq_wheel** with **wheelS_state_eq**, for the original fill-form wheel.*

The run-dependent level is computable from the pass trace. At $L = 1$, **wheelS_unit_balanced** and **runW_unit_balanced** reduce the band to the indivisible-claims balancedness condition $a_i \leq a_j + 1$ for every unexhausted j , matching the benchmark of [Herrero and Martínez \(2008\)](#).

Corollary 6.5 (Vanishing-lot analytic limit). *Fix a participant count $m \geq 1$. For run k , let $Q^{(k)} \in \mathbb{N}^m$, lot L_k , fill vector $a^{(k)}$, and executable water level w_k be the quantities in Theorem 6.4. For any real scale $c_k > 0$,*

$$\max_{1 \leq i \leq m} \left| \frac{a_i^{(k)}}{c_k} - \min \left\{ \frac{Q_i^{(k)}}{c_k}, \frac{w_k}{c_k} \right\} \right| \leq \frac{L_k}{c_k}.$$

Consequently, if $L_k/c_k \rightarrow 0$, the normalized wheel allocations converge uniformly to their constrained-equal-awards profiles. If also $Q_i^{(k)}/c_k \rightarrow q_i$ for every i and $w_k/c_k \rightarrow s$, then $a_i^{(k)}/c_k \rightarrow \min\{q_i, s\}$ for every participant.

Proof. The machine-checked finite inequality in Theorem 6.4 gives $|a_i^{(k)} - \min\{Q_i^{(k)}, w_k\}| \leq L_k$ for each i . Embed the natural numbers in $\mathbb{R}$, divide by c_k , and take the maximum over the fixed finite set of participants. The first limit follows by squeezing. The second follows from continuity of $\min$. $\square$

This real-valued limit is an analytic corollary of the checked integer bound, not an additional Lean declaration. It identifies the precise asymptotic regime: the lot must vanish relative to the economic scale, not necessarily in absolute shares.

Theorem 6.6 (Impossibility; **serial_band_incompatible**, **ex_incompatible**). *Let two participants hold $q_1 \geq x$ and $q_2 > 0$ with $L < x$. Every allocation satisfying Axioms 5.1 and 5.2 is $(x, 0)$, which violates Axiom 5.3 for every water level. Instance: sizes $(200, 200)$, $L = 100$, $x = 200$.*

Within these definitions, seriality and the band condition are jointly unsatisfiable on the stated family of two-order instances. A hybrid can apply serial priority to a setter tranche and a banded allocation to a separate residual pool; the theorem does not classify every priority feature of an actual venue.

Theorem 6.7 (Ladder uniqueness; **LadderSerial**, **ptLadder**, **ladder_unique**). *Extend seriality across levels: a deeper level filled implies the shallower level filled in full. Any ladder allocation that is level-wise feasible and serial, and serial across levels, equals the price-time ladder.*

Remark 6.8 (Class versus point). Axiom 5.3 does not determine the allocation: the pointer is free, and distinct pointers can yield identical fills (Section 11). The band therefore defines a compatibility class and a proved property of the modeled wheel, not a characterization of venue parity. By contrast, feasibility plus seriality characterize the modeled price-time output.

7 Structure theorems

Theorem 7.1 (Conservation and caps; `wheelRound_sum`, `wheelRound_cap`, `wheelRound_le_L`, `wheel_sum`, `wheel_cap`, `rule737_sum`, `rule737full_sum`, `priceTime_cap`, `priceTime_sum_le`). *At every tier, fills respect sizes, a wheel visit gives at most one round lot, and fills plus leftover equal the incoming quantity.*

Theorem 7.2 (Displayed precedence; `t3_only_after_t2`). *For the specific three-tier operator `rule737full`, non-displayed interest receives quantity only after the displayed tier leaves positive quantity. The result concerns this three-tier operator, not every venue’s allocation rule.*

Theorem 7.3 (Completeness; `tot`, `passS_tot_lt`, `wheelS_complete`, `wheelS_exhausts`, `wheel_exhausts`). *With $L \geq 1$, an un-truncated pass over a level with positive total remaining strictly reduces that total. Hence if quantity is left over, either the level is exhausted or fewer than `tot` passes were available; with fuel `tot + 1`, positive leftover forces exhaustion, and the fills then equal the displayed sizes.*

At this canonical fuel bound, positive leftover certifies exhaustion of the level. With a smaller fuel bound it can instead reflect computational truncation.

Theorem 7.4 (Price priority on a ladder; `matchAsks`, `matchAsks_conserve`, `price_priority`, `no_fill_above_limit`). *A buy with limit ℓ walks the ask ladder over levels priced $\leq \ell$, price–time within each. Fills plus leftover equal x ; a level above the limit receives nothing; and if any deeper level receives quantity, every order at the top marketable level is filled in full.*

Theorem 7.5 (Price priority over parity; `matchNYSE`, `matchNYSE_conserve`, `nyse_price_priority`). *On a ladder of hybrid levels (price, setter tranche, parity sizes), with $L \geq 1$ and fuel $\sum \text{parity} + 1$ at the top marketable level: if any deeper level receives quantity, the top level’s setter tranche was filled in full and its parity fills sum to its displayed size.*

Price priority is seriality one level up. A serial venue is price–time on the ladder, each level price–time on its queue; a hybrid venue is price–time on the ladder, each level setter–then–parity on its queue.

Example 7.6 (`ex_book`, `ex_nyse`). Serial ladder $\{10.00:(200, 300), 10.01:(400), 10.03:(500)\}$: a 700-share buy with limit 10.01 fills (200, 300) then 200; with limit 10.00 it fills 500 and leaves 200. Hybrid ladder $\{10.00: \text{setter } 200, (300, 500, 0, 1000)\}, \{10.01:(400)\}$: a 2200-share buy exhausts the first level ($200 + 1800$) and takes 200 at 10.01; a 1500-share buy stops inside the first level with wheel fills (300, 500, 0, 500).

Proposition 7.7 (Integer wheels need not preserve depth-share domination; `wheel_book_share_counterexample`). *Let the book be first, with sizes (1000, 100), lot $L = 100$, and incoming quantity $x = 100$. Both the pass wheel and the persistent pointer wheel give (100, 0). The book receives all executions while supplying only 1000/1100 of the depth, and $100 \cdot 1100 > 1000 \cdot 100$. Thus the benchmark’s multiplicative share inequality does not transfer unchanged to either integer wheel.*

Lemma 7.8 (Book share under water-filling; `book_share_le_depth_share_mem`). *Under water-filling at level s , if a participant of size Q_b is uncapped ($s \leq Q_b$) and at least as large as every other ($q_i \leq Q_b$), then $\min\{Q_b, s\} \cdot \sum_i q_i \leq Q_b \cdot \sum_i \min\{q_i, s\}$: its fill relative to the others’ is at most its depth relative to theirs.*

Here the book is one of the participants, and its being largest and uncapped is an explicit hypothesis. The lemma concerns the water-filling benchmark; the preceding proposition isolates exactly why the finite-increment wheels differ.

8 Allocation as rationing, and the pointer theorem

8.1 The identification

A claims problem is a vector of claims $c \in \mathbb{N}^m$ and an amount $E \leq \sum c$ to divide; a rule assigns awards $a \leq c$ summing to E (Young, 1987, 1994; Thomson, 2003). A queue at a price with incoming quantity x is a claims problem with claims q and amount $\min\{x, \sum q\}$. Under this identification:

- price-time is the *sequential priority* rule for the arrival order (Moulin, 2000);
- constrained equal awards (called water-filling in the executable kernel) assigns $\min\{q_i, s\}$ (Herrero and Villar, 2001);
- the round-lot wheel is a finite-increment mechanism with a constrained-equal-awards approximation guarantee and a rotation order, related to indivisible claims allocation (Herrero and Martínez, 2008).

Theorem 6.2 is then a characterization of sequential priority by seriality, Theorem 6.4 a finite-increment approximation theorem for constrained equal awards, and Theorem 6.6 an incompatibility result for the stated priority and balance axioms. The rationing literature supplies a further composition requirement.

Axiom 8.1 (Composition; Moulin, 2000). Allocating x , then allocating y against the residual claims, gives the same awards as allocating $x + y$ at once.

Composition is what makes an allocation rule well defined on a *stream* of executions rather than on a single one. An exchange processes incoming orders one at a time; without composition, the fills a resting order receives would depend on how the contra-side quantity happened to be chopped into orders.

Definition 8.2 (Stateful stream consistency). For a transition system, the analogous property compares final states: running x and then y should agree with running $x + y$, subject to the termination conditions of the executable transition system. This is distinct from Axiom 8.1, whose domain contains only residual claims.

8.2 What composes, and on which index

Theorem 8.3 (Estate composition and level composition). Lean: `priceTime_comp`, `waterFill_level_comp`, and `waterFill_two_equal_no_single_share`. *Price-time satisfies estate composition (Axiom 8.1):*

$$\text{priceTime } Q (x + y) = \text{priceTime } Q x \oplus \text{priceTime } (Q - \text{priceTime } Q x) y.$$

The benchmark instead satisfies a water-level identity, $\min\{q, s + t\} = \min\{q, s\} + \min\{q - \min\{q, s\}, t\}$ pointwise. This identity is not Axiom 8.1 on an integer estate domain: for claims (10, 10), every natural water level clears an even number of shares, so no level represents the one-share estate. Thus natural water levels do not parameterize the full integer estate domain $0 \leq E \leq \sum q_i$.

Theorem 8.4 (Estate-indexed completion and the reset obstruction; `discreteCEA_feasible`, `discreteCEA_reset_not_composable`). Define `discreteCEA` Q E by running the unit-increment wheel on claims Q and estate E , with a claims-determined fuel bound. For every Q and E , this

allocation is feasible: it respects every claim and awards exactly $\min\{E, \sum_i q_i\}$. It therefore completes the integer estate domain omitted by natural water levels. But if its deterministic rotation is reset for each estate, it fails Axiom 8.1: on $Q = (10, 10)$, one unit followed by one unit yields $(2, 0)$, whereas the two-unit estate yields $(1, 1)$.

A deterministic tie-break that restarts from a fixed participant is an exact indivisible allocation rule for each estate, yet not a stream-consistent rule on residual claims. Persisting its rotation produces the stateful wheel below and restores conditional stream consistency; it does not turn the stateful theorem into the stateless Axiom 8.1.

8.3 The pass-reset wheel does not

Theorem 8.5 (*wheel_not_path_independent*). *Two participants of 300, $L = 100$. Allocating 100 and then 100, restarting the pass at the pointer each time, gives $(200, 0)$; allocating 200 at once gives $(100, 100)$.*

The pass-reset wheel begins each execution with a fresh pass at the first participant. It is banded (Theorem 6.3) and conservative, but it is not path independent, and a resting order's fill under it depends on the chopping of the contra flow.

8.4 The pointer-and-lot wheel does

Let the state carry, in addition to each participant's (allocated, remaining), the rotation itself—the internal list is kept rotated so that its head is the pointer—and the round-lot allowance still owed to the head (**WState**). A visit (**stepW**) rotates the dead prefix to the back (**findLive**), gives the live head $\min\{\text{lot}, \min\{r, x\}\}$, and either advances the pointer and resets the allowance to L (when the visit completed the lot or exhausted the participant) or stays with the allowance reduced (when the incoming quantity ran out mid-lot). A run (**runW**) iterates visits until quantity or remaining size is exhausted. Participant identifiers are preserved through rotation; output comparisons are made after restoring canonical label order.

Theorem 8.6 (Balance of the pointer wheel; **AllAtP**, **AllAtP_band**, **stepW_allAtP**, **runW_band**, **runW_balanced**). *For $L > 0$, every state returned by a fuel-bounded **runW** execution initialized with zero allocations satisfies the band predicate for some w : all allocations are at most $w + L$, and every unexhausted participant has allocation at least w . Exhausted participants may lie below w . Consequently, no participant is more than one round lot ahead of an unexhausted participant. This is the same **runW** operator used in Theorem 8.8.*

Proof idea. The pointer-relative invariant **AllAtP** records a low-then-high cycle boundary. The current live head has allocation plus unused lot allowance exactly $w + L$; participants not yet visited in the cycle remain at w , participants already visited are at $w + L$, and exhausted records never exceed that ceiling. A partial visit reduces the allowance by exactly the fill, a completed visit moves the head across the boundary, and completing a cycle raises w by L . Thus every visit preserves the invariant, which implies the band. $\square$

Lemma 8.7 (Step composition; **stepW_split**). *Against a non-exhausted state, either the first visit with x and with $x + y$ coincide, or the visit with x is partial and the visit with $x + y$ equals the visit with y from the partial state, with fills adding.*

Theorem 8.8 (Path independence of the pointer wheel; `runW_stable`, `runW_comp`). *Let $r_1 = \text{runW } L \ n_1 \ st \ x$ and $r_2 = \text{runW } L \ n_2 \ r_1.1 \ y$. If $r_1.2 = 0$ and `Stopped` $r_2.1 \ r_2.2$, then*

$$\text{runW } L \ (n_1 + n_2) \ st \ (x + y) = r_2.$$

Thus a first run that fully allocates x , followed by a terminating run of y , agrees exactly with the combined run at the summed fuel. The theorem is uniform in the round lot, participant count, pointer position, and partially consumed current lot; it does not remove its completion, termination, or fuel hypotheses.

Proof idea. Induction on the fuel of the first run. Each visit is either identical in the split and joint runs (the lemma’s first case), in which case the induction hypothesis applies to the residual, or partial (the second case), in which case the first run has ended with quantity zero, the joint run’s next visit is exactly the second run’s first visit, and the remaining visits agree by the stability of terminated runs under extra fuel. $\square$

Corollary 8.9 (Canonical stream consistency on reachable states; `runWFull`, `runWFull_comp_reachable`). *Let st be any state returned by `runW` from `initW`, and let `runWFull` choose fuel $\text{totW}(st.S) + 1$. If $L > 0$ and the first estate satisfies $x \leq \text{totW}(st.S)$, then for every y ,*

$$\text{runWFull } L \ st \ (x + y) = \text{runWFull } L \ (\text{runWFull } L \ st \ x).1 \ y.$$

The reachable-state allowance invariant makes positivity automatic. The supporting descent and conservation lemmas prove that canonical fuel terminates and that the first admissible estate clears. Allowance positivity, completion, and stopping therefore follow from reachability, canonical fuel, and the first-estate bound. The lower-level theorem `runWFull_comp` covers arbitrary supplied states and therefore retains a positive-current-allowance premise to exclude malformed states with allowance zero.

Example 8.10 (`pointer_wheel_instance`, `pointer_wheel_odd_instance`). On the counterexample above, the persistent wheel allocates 100 then 100 as (100, 100), the same as 200 at once; 150 then 150 gives the same state as 300 at once, the partial lot carrying across the boundary.

8.5 What the theorem establishes

Exchange documentation specifies that parity allocation proceeds from an allocation wheel and that the wheel advances across executions (New York Stock Exchange, 2024). Theorems 8.6 and 8.9 place balance and exact split/combined equality on one executable construction. Example 11.4, Theorem 11.6, and Theorem 11.7 below use explicit, different observation contracts: the four-party collision requires pointer information beyond queue and fills; the fiber theorem gives exact pointer necessity and sufficiency conditional on a common authenticated observation; and the mL result bounds a complete predictive summary without separate side inputs. Corollary 11.8 proves that cumulative allocation determines the allowance on every reachable live-head state. These statements concern the modeled wheel and do not characterize all banded stream-consistent mechanisms.

Plain-language reading. If a 200-share execution is processed as two consecutive 100-share executions, a stream-consistent mechanism should end in the same state either way. A wheel that forgets its position can fail. The persistent wheel passes this test for every reachable state and admissible first estate. Its unused allowance can always be recovered from the cumulative fill at the current pointer; the pointer label remains the phase datum that aligned fills do not necessarily reveal.

9 Cross-market structure

9.1 Consolidation

The definitions in this section are executable analytic kernels for selected Regulation NMS concepts. The integrated regulatory specification in Section 14 supplies the definitions, exceptions, procedural duties, and versioned parameters needed for clause traceability. Unless stated otherwise, numerical parameters describe the operational compliance regime on September 5, 2026 (U.S. Securities and Exchange Commission, 2005, 2025; Office of the Federal Register and U.S. Government Publishing Office, 2026a).

Definition 9.1 (`nbo`, `nboSH`). The model NBO is the minimum protected offer across the supplied venue-top quotations. A quote is protected only when it is automated, at the venue top, and at least the supplied round-lot size L .¹ Under the modeled self-help exception, the affected venue’s quotations are excluded.

Theorem 9.2 (`nbo_le_protected`, `oddlot_not_protected`, `selfhelp_raises_nbo`). *Every protected offer is at or above the NBO; an odd-lot or manual quote is never protected and never sets it; removing a venue’s quotes can only raise it.*

Example 9.3 (`ex_nbo`). With $L = 100$, venue v_1 supplies an automated 300-share offer at 10.02; v_2 an automated 50-share offer at 10.01; v_3 a manual 200-share offer at 10.01; and v_4 an automated 500-share offer at 10.03. The model NBO is 10.02.

9.2 Order protection

Definition 9.4 (`tradeThroughBuy`, `sweepThenExecute`). In the model, a buy executed above the NBO is a trade-through. The sweep operator sends an intermarket sweep to each supplied protected quotation better than execution price px , filling $\min\{\text{size}, x\}$ in list order, then records the remainder at px .

Theorem 9.5 (`compliant_of_le`, `route_to_nbo_no_trade_through`, `sweep_conserve`, `iso_compliance`). *Execution at or below the NBO is compliant. Sweep fills plus quantity executed at px equal x , and if any quantity is executed at px then every better protected quote received a sweep for exactly its size.*

Example 9.6 (`ex_sweep`). On the quotes above, a 1000-share sweep at 10.03 sends 300 to v_1 and executes 700 locally.

The local remainder is not capped by modeled venue liquidity. The code also omits the one-second window, flickering-quotation treatment, quotation availability and timing, and Rule 611 exceptions other than the modeled ISO and self-help cases. Accordingly, `iso_compliance` is a theorem about this abstraction rather than certification of full Rule 611 compliance.

¹The parameter L is security-specific, not a universal 100-share constant. Rule 600(b)(93) assigns 100 shares when the prior evaluation-period average closing price is at most 250, 40 shares from 250.01 through 1,000, 10 shares from 1,000.01 through 10,000, and one share above 10,000; a new NMS stock receives 100 shares. The compliance date for this round-lot definition was November 3, 2025. The October 2025 relief concerned exchanges’ conforming rule-change filings, not postponement of the tiered definition itself. The examples use $L = 100$, while the theorems are uniform in L (Office of the Federal Register and U.S. Government Publishing Office, 2026a; U.S. Securities and Exchange Commission, 2024a, 2025).

9.3 Tick, lock, fee

Definition 9.7 (`rule612`, `notLockedOrCrossed`, `slideBuy`, `feeCap`, `netCost`). Under the legacy parameterization encoded here, quotes at or above \$1 are in whole cents and the access-fee cap is \$0.003 per share. The NBB is strictly below the NBO; a buy that would lock is displayed one modeled tick below the NBO; net cost is price plus fee.

Theorem 9.8 (`effective_tick`, `relative_tick_antitone`, `slide_no_lock`, `slide_le_limit`, `netCost_le_cap`, `fee_adjusted_spread`). *Tick compliance plus no lock or cross forces a quoted spread of at least one cent at or above \$1; the relative tick is non-increasing in price; sliding never locks and never exceeds the limit; the fee-adjusted spread exceeds the quoted spread by at most two fee caps.*

The SEC adopted a conditional \$0.005 minimum increment and a \$0.001 access-fee cap in 2024 (U.S. Securities and Exchange Commission, 2024a). Relief granted in 2025 was extended on June 11, 2026: compliance with amended Rules 610(c) and 612, and with Rule 600(b)(89)(i)(F), is postponed until the first business day of November 2027 (U.S. Securities and Exchange Commission, 2025, 2026b). Thus the analytic kernels `rule612` and `feeCap` encode the regime operative on the paper’s September 2026 date, while the clause layer separately encodes both legacy-relief and delayed amended functions in `minimumTick612` and `accessFeeCap610`. The compatibility theorems `rule612_iff_minimumTick612_legacy` and `feeCap_eq_accessFeeCap610_legacy_dollar` prove that the legacy clause branches reduce to the analytic kernels. A separate June 2026 proposal to rescind Rule 611 and Rule 610(e) remained a proposal on the paper’s stated date (U.S. Securities and Exchange Commission, 2026a).

9.4 Two market-data clocks

Theorem 9.9 (`sip_ge_venue`, `sip_order_preserved`, `sip_order_inversion_exists`, `sip_inversion_bound`). *Consolidated time is never earlier than venue time; equal latencies preserve order; inversions exist; an inversion requires the latency difference to exceed the venue-time gap.*

9.5 Order types

Theorem 9.10 (Selected order-type properties). Lean: `midpoint_inside`, `midpoint2_strict_inside`, `mpl_no_trade_through`, `fillReserve_conserve`, `fillReserve_replenish`, `dOrder_within`, `dOrder_compliant`, `marketPeg_no_lock`, `primaryPeg_mono`, `postOnly_rests_iff`, and `postOnly_never_takes`. *The grid-valued midpoint $\lfloor (NBB + NBO)/2 \rfloor$ lies in $[NBB, NBO)$ when $NBB < NBO$; it need not be strictly above the NBB for a one-grid-unit spread. In doubled-price units, the exact midpoint $NBB + NBO$ is strictly inside every non-locked spread. The floored kernel never trades through. A reserve order’s display plus reserve falls by exactly the fill, and replenishment occurs only when the display was exhausted, to at most the original display. A discretionary order executes within $[display, limit]$ and is protection-compliant when its limit is at or below the NBO. A market peg with $0 < offset \leq NBO$ never locks; a primary peg re-prices monotonically with the NBB. A post-only order rests only strictly below the best contra quote and never removes liquidity.*

These are one-step semantic kernels. In particular, `fillReserve` performs at most one refresh. Repeated matching, queue re-entry, repricing, rejection, routing, and attribute compatibility are represented as explicit certification duties in `ClauseComplete.lean`; they are not implied by the one-step arithmetic theorem. Nasdaq’s priority rules distinguish displayed and non-displayed interest beyond the within-queue price–time primitive modeled here (The Nasdaq Stock Market LLC, 2026).

9.6 Bands and tests

Theorem 9.11 (`luldBands`, `lower_le_ref`, `ref_le_upper`, `bands_widen`, `limitState_boundary`, `shortSaleOK`, `no_short_at_or_below_nbb`, `effSpread2`, `eff_le_quoted`, `effSpread2_neg_iff`, `midpoint_zero_eff`). *Price bands $\text{ref} \pm \lfloor \text{ref} \cdot \text{pct}/100 \rfloor$ contain the reference and widen in the percentage; a limit state has one side of the market on a band. With the short-sale circuit breaker in effect a short sale executes only strictly above the NBB. The signed buy-side kernel $\text{effSpread2 } px \text{ nbb } nbo = 2px - (nbb + nbo)$ is evaluated in $\mathbb{Z}$: it is at most the quoted spread when $px \leq nbo$, vanishes at an exact grid midpoint, and is negative exactly below the midpoint. It preserves below-midpoint magnitude; a complete Rule 605 statistic still requires side, order-class, time, and aggregation inputs.*

The clause layer represents the LULD obligations in a record for reference recomputation, tiering, late-day changes, rounding, limit and straddle states, pauses, reopening, dissemination, policies, and operative overnight protection. It derives the Rule 201 ten-percent trigger, enumerates every short-exempt reason and its internal conditions, and represents the complete amended Rule 605 order-class and disclosure-column schema. The kernel `effSpread2` alone is not a report generator; `rule605Complies` checks whether a supplied report record contains every required component.

10 Auction structure

Definition 10.1 (`demandAt`, `supplyAt`, `matchedAt`, `imbalanceAt`, `crossKey3`, `crossPrice3`, `dmmFacilitate`). Orders are (limit, size). Demand at p is the size of bids with limit $\geq p$; supply at p the size of asks with limit $\leq p$; matched volume is their minimum; imbalance is the pair of one-sided excesses. The clearing key encodes, lexicographically, matched volume, then $B - \text{imbalance}$, then $D - |p - \text{ref}|$, for bounds B, D ; the auction price maximizes the key over candidates. The designated market maker absorbs residual imbalance up to a committed size.

This is the proved generic lexicographic auction kernel. Clause-level NYSE and Nasdaq records separately require the full definitions, eligible-interest treatment, entry and cancellation windows, imbalance messages and schedules, collars, priority, uniform price, official-price publication, LULD and halt variants, DMM and floor responsibilities, and contingency procedures. Facilitation remains a two-number theorem; satisfaction of the wider procedural record requires external evidence.

Theorem 10.2 (`demandAt_antitone`, `supplyAt_mono`, `crossPrice_maximizes`, `crossPrice_global`, `crossPrice2_lex`, `crossPrice3_lex`, `dmm_clears`, `dmm_conserve`). *Demand is antitone and supply monotone in price. Over a nonempty finite candidate list the clearing price maximizes matched volume; over the ask prices it is a global maximizer; with the two- and three-level keys, assuming B bounds every candidate imbalance and D bounds every candidate reference distance, it also minimizes imbalance and then reference distance among maximizers. Facilitation clears the residual when the commitment covers it and conserves quantity.*

Example 10.3 (`ex_cross`, `ex_auction`). Bids $\{(10.02, 500), (10.01, 300), (10.00, 400), (9.99, 200)\}$, asks $\{(9.98, 300), (10.00, 300), (10.01, 400), (10.03, 500)\}$: the auction clears at 10.01 with 800 matched and a 200-share sell imbalance (10.00 would match 600, 10.02 only 500); with reference 10.00 the three-criterion price is still 10.01; a 200-share residual against a 500-share commitment clears.

Plain-language reading. An auction tests candidate prices. At each price it computes how many shares buyers and sellers could trade, selects the price with the greatest matched volume, then uses imbalance and distance from a reference price as tie-breakers. “Lexicographic” means applying those criteria in that fixed order, like sorting first by surname and then by given name.

11 Observability structure

Theorem 11.1 (Identification). Lean: `parity_count_fingerprint`, `hidden_depth_lower_bound`, and `reserve_certificate`. For an observer who knows L and can delimit one un-truncated parity pass in a participant-aligned allocation record, if every participant holds at least a round lot then the pass fills are $(L, \dots, L)$ of length m : its length identifies the participant count. A public trade tape ordinarily lacks that participant alignment; without authenticated alignment, the known lot size, and the pass boundary, the theorem makes no surveillance claim. Separately, an execution of x that clears a level certifies non-displayed depth of at least $x - Q^s - \sum r_i$.

Theorem 11.2 (Invisibility). Lean: `priceTime_zero`, `priceTime_tail_invisible`, and `priceTime_boundary_invisible`.

If $x \leq q_1$, the queue-aligned price-time allocation vector is the same for every queue tail of the same length. More strongly, any two head sizes at least x , followed by equally long tails, generate the same vector. Thus aligned fills identify sizes only for a fully exhausted prefix; they give a lower bound for a partially filled marginal order and do not identify the remaining tail.

Theorem 11.3 (No parity conformance from queue and aligned fills alone). Lean: `parity_pointer_not_identified_from_prints` and `no_parity_conformance_from_queue_and_prints`.

For the participant-labelled queue (300, 500, 200, 1000), two distinct pointer states produce the same 1,100-share queue-aligned fill vector (300, 300, 200, 300) but reach distinct terminal states; the next 100-share allocation goes to different participants. Moreover, one identical candidate vector for the same queue conforms under one pointer and fails under the other. Therefore no Boolean function of the queue and aligned fills alone can be correct on both states: a parity conformance check must receive the pointer state or equivalent authenticated evidence.

Example 11.4 (Four-party future-separation witness; `FutureSufficient4`, `future_sufficient_summary_requires_pointer_information`). Let S map wheel states to an arbitrary summary type. Suppose there is a predictor that, for every four-party state and next quantity, correctly returns the next queue-aligned fill vector from only the current canonical queue, current aligned cumulative fills, $S(st)$, and that quantity. Then S must assign different values to `pointerA` and `pointerC`. In particular, no universally future-sufficient summary can collapse the pointer distinction erased by queue and fills.

Proof idea. The two witness states have the same canonical queue and zero cumulative fills. If their summaries were also equal, the predictor would receive identical inputs for the next 100 shares and would have to return the same vector. Evaluation proves that the actual next vectors differ, a contradiction. $\square$

Theorem 11.5 (General nondegenerate phase separation; `ValidWheelPhase`, `validWheelPhase_probe_equiv_implies`). Let st and su each have a live head, a distinct live successor, a current allowance in $\{1, \dots, L\}$, and more than L shares remaining at every participant. Write their head pointer-allowance pairs as (p, a) and (q, b) . If their participant-labelled new-fill vectors agree for every estate $x \leq L$, then $p = q$ and $a = b$.

Proof idea. The one-share probe identifies the head pointer. Once the pointers agree, suppose without loss that $a < b$. At estate $a + 1$, the first state completes its allowance and gives the next share to its distinct successor, so p receives a ; the second state gives all $a + 1$ shares to p . This contradicts probe equivalence. The hypotheses keep both probes away from exhaustion boundaries. The result is a necessity statement across arbitrary accumulated allocations, successor labels, and

tails in this nondegenerate domain; equality of pointer and allowance alone does not erase differences in those other state components. $\square$

Fix an arbitrary authenticated observation O and an indexed family $\{st_p : 0 \leq p < m\}$. A summary S is *future sufficient on this fiber* if one decoder, given $O(st_p)$, $S(st_p)$, a future estate x , and a participant label i , returns the actual canonical next fill $F_L(st_p, x)_i$ for every p, x, i (**WheelFiberSufficient**).

Theorem 11.6 (Fiberwise pointer minimality; **wheelFiberSufficient_pointer_injective**, **wheelPointerCode_fiber_sufficient**). *Let $L > 0$. Suppose every st_p is a nondegenerate wheel phase whose live head label is p , and $O(st_p)$ is the same for all $p < m$. Then every future-sufficient summary S is injective on the family:*

$$S(st_p) = S(st_q) \implies p = q.$$

Conversely, the head label p itself is a future-sufficient summary for the fixed indexed family. Thus the pointer is an exact minimal predictive code on this common-observation fiber: if all m indices occur, precisely m summary values suffice and fewer cannot.

Proof idea. If two summary values were equal, their observation inputs would also be equal, so the common decoder would make the two states agree on every future probe. Theorem 11.5 then forces their head labels to agree. For sufficiency, the decoder uses the supplied head label to select the fixed family member and executes **nextWFill**. The necessity argument permits arbitrary claims, accumulated fills, successors, tails, observation types, summary types, and decoders satisfying the stated contracts; the converse is relative to the fixed one-state-per-pointer family. $\square$

A finite family of executable phase states. Fix $m \geq 2$ and $L > 0$. Label the participants $0, \dots, m-1$ in cyclic order and give each original claim $2L+1$. For $0 \leq p < m$ and $1 \leq a \leq L$, define $P_{p,a} = \text{phaseState } mLpa$ by putting p at the head, with record $(p, L-a, L+1+a)$ and current allowance a . Every other label follows in cyclic order with record $(i, 0, 2L+1)$. The declarations **phaseState_length**, **phaseState_labels_unique**, **phaseState_capacities**, and **phaseState_claims_large** verify the participant count, distinct labels, common original capacities, and remaining claims strictly above L . Moreover, **phaseState_reachable** proves that $P_{p,a}$ is reached by initializing that cyclic queue at p and running $L-a$ shares with **runWFull**.

Write $F_L(st, x)_i$ for the new allocation to label i produced by **runWFull**; the executable definition **nextWFill** subtracts that label's cumulative allocation before the run from its allocation afterward. Define $st \sim_L su$ when $F_L(st, x)_i = F_L(su, x)_i$ for every x and label i (**WheelFutureEquivalent**). A complete predictive summary has a decoder for those fills from the summary and (x, i) alone, with m, L fixed; it receives no separate queue or cumulative-fill input (**WheelPhaseSufficient**).

Theorem 11.7 (Finite executable phase identification and minimal predictive summaries). **Lean:** **phaseState_probe_equiv_iff**, **phaseState_future_equiv_iff**, **wheelPhaseSufficient_injective**, and **wheelPhaseCode_sufficient**.

For the valid phases above,

$$P_{p,a} \sim_L P_{q,b} \iff (p, a) = (q, b).$$

It is enough to compare estates $x \leq L$. Every complete predictive summary is injective on these mL reachable states, and the code (p, a) is sufficient. Hence the family has exactly mL behavioral classes and any fixed-length binary encoding of a complete predictive summary needs at least $\lceil \log_2(mL) \rceil$ bits.

Proof. The actual canonical wheel gives its first share to p (`phaseState_first_fill`), so different pointers are separated by $x = 1$. For equal pointers and $a < b$, take $x = a + 1 \leq L$. In $P_{p,a}$, the first a shares finish the current allowance and the successor receives one share; p receives a (`phaseState_allowance_probe`). In $P_{p,b}$, all $a + 1$ shares go to p (`nextWFill_head`). These conclusions use `runWFull` and its composition theorem directly. The case $b < a$ is symmetric. Equal phases reconstruct equal finite states, so all subsequent executions agree. A decoder cannot give the same code to two states separated by a probe; conversely `wheelPhaseCode` records the head label and allowance and reconstructs the family member. Lean checks reachability, the separating executions, equivalence, injectivity, and sufficiency. Counting the m pointer values and L allowance values gives the cardinality; $2^b \geq mL$ gives the stated binary-encoding corollary. $\square$

This uses the future-distinguishability principle underlying automaton minimization (Nerode, 1958), applied to a specified reachable family of finite `WState` values. Equivalence on every finite quantity stream gives the same classes: a single-estate stream supplies each separating probe, while equal phases give identical initial states and therefore identical streams. It does not characterize every banded mechanism or minimize all raw `WState` representations. One-live-participant states, exhausted claims, and dead-prefix representations can erase distinctions elsewhere.

Corollary 11.8 (Reachable allowance-phase identity; `runW_allowance_phase`, `allowance_from_cumulative`). *Fix $L > 0$ and any execution of `runW` from `initW`. If the returned state has a live head p with cumulative allocation A_p , then its current allowance is*

$$a = L - (A_p \bmod L).$$

Thus knowing the pointer and its aligned cumulative allocation determines the allowance on every reachable live-head state; no cycle-phase premise is supplied by the caller. The supporting invariant `runW_atWheelPhase` proves $A_p + a = kL + L$ for some k , and `allowance_from_cumulative` discharges the arithmetic identity.

Corollary 11.9 (Exact pointer state with aligned cumulative fills; `phaseState_full_observation`, `phaseState_full_summary_injective`, `phaseState_full_pointer_sufficient`). *For $m \geq 2$ and $L > 0$, restrict to the reachable full-allowance states $\{P_{p,L} : 0 \leq p < m\}$. Their original capacities are all $2L + 1$, and their participant-aligned cumulative-allocation observation is the identical vector $(0, \dots, 0)$. Every future-sufficient extra summary must distinguish all m pointers, and the pointer label itself is sufficient. The exact extra-state count on this authenticated-observation fiber is therefore m ; any fixed-length binary encoding needs at least $\lceil \log_2 m \rceil$ bits.*

The identity in Corollary 11.8 explains why the observation contracts must stay distinct. Theorem 11.7 bounds a complete summary that must carry all predictive information for its family, so its mL count is not an extra-bit bound after cumulative fills are supplied. Theorem 11.6 and Corollary 11.9 instead prove the exact pointer requirement on common-observation fibers, while allowance recovery establishes redundancy of the separate allowance field on the larger class of all reachable live-head states. A global minimal quotient across arbitrary claims and degeneracies is not asserted.

Example 11.10 (A finite allowance probe; `ex_phase_state_separation`). For $m = 3$, $L = 100$, and pointer $p = 1$, an estate of 41 gives label 1 exactly 40 shares in $P_{1,40}$ but 41 in $P_{1,80}$. Both states are reachable and all remaining claims exceed the probe. The different cumulative allocations (60 and 20 at the pointer) also reveal why providing those allocations changes the information question. Claims above L are needed only to keep this finite separating probe away from exhaustion; no infinite-backlog assumption is made.

For comparison, `saturatedOwner` defines an idealized infinite-backlog phase formula with elementary trace-injectivity lemmas. Those lemmas are separate from, and unnecessary for, the finite executable theorem above.

Plain-language reading. A public trade tape records executions but does not normally expose the full prior order queue or a participant-aligned allocation vector. Even when an auditor has the stronger aligned record used by these theorems, different hidden queues or pointer states can generate the same fills. The results state exactly what can still be inferred and what cannot be reconstructed from those inputs alone.

12 Conformance

Corollary 12.1 (Decidable, complete, and sequentially composable; `conformsPT`, `conformsPT_iff`, `conformsPT_comp`). *The Boolean “aligned fills equal the engine output” is decidable and equivalent to “aligned fills satisfy Axioms 5.1 and 5.2.”*

A surveillance rule is, in this foundation, a decidable predicate with a completeness theorem; a flagged allocation record names the axiom it violates. If one aligned fill vector conforms to the original queue and a second conforms to the residual queue, `conformsPT_comp` proves that their pointwise sum $A \oplus B$ conforms to one execution of the combined quantity. Acceptance is therefore stable when two individually conforming executions are aggregated. `AllocationRecord.lean` implements the one-execution predicate as a program reading lines `queue;x;fills` and reporting OK or the violated property with the first offending index (`CAP`, `CONSERVATION`, `SERIAL`).

`Conformance.lean` exercises the checker on synthetic allocation records. A fixed-size harness accepts 200/200 engine-generated records and rejects 200/200 single-transfer mutations. A second harness varies incoming size, accepts 200/200 honest records, and rejects 800/800 mutations spanning forward transfer, backward transfer, over-fill, and under-fill. These are deterministic regression tests, not statistical estimates and not an independent implementation of the axioms. The command-line parser rejects any malformed numeric token rather than deleting it. An additional executable regression script checks 12 cases covering valid input, parse rejection, and named length, cap, conservation, and seriality failures. `Collisions.lean` displays explicit distinct queues that generate identical price–time fill vectors; the examples are witnesses, not a population estimate from an unspecified sampling design.

Plain-language reading. The checker receives the known or reconstructed queue, incoming quantity, and a privileged queue-aligned fill vector. It either returns OK or names the first violated rule. Passing means the record matches the formal price–time model for those inputs. It does not prove that the supplied queue or order-to-fill alignment is authentic, which is why consolidated trade data by itself is insufficient.

13 Rule-by-rule coverage register

This compact register accounts for every analytic market-rule family encoded in the Lean package. The kernel column names the executable abstraction, the guarantee column states only what Lean proves about it, and the boundary column identifies what that theorem leaves outside. The last column points to the corresponding clause-layer record where one exists. Analytic-only entries have no corresponding clause-layer certificate.

Family	Kernel	Proved guarantee	Boundary	Clause-layer object
Price-time	<code>priceTime;</code> <code>matchAsks</code>	Unique feasible serial allocation; caps, conservation, monotonicity, and price priority.	Queue construction, modifications, cancellations, hidden priority, self-trade prevention, and venue tie-breaks are inputs or omitted.	VenueBookState
Parity	<code>wheel;</code> <code>discreteCEA;</code> <code>runWFull</code>	Estate rule: feasibility but reset-composition failure. Persistent wheel: one-lot band, reachable composition, allowance recovery, general fiberwise pointer minimality, and exact finite-family quotients.	Global raw-state quotient can collapse at claim boundaries or with one live participant; eligibility, aggregation, label restoration, and within-participant ranking remain inputs (New York Stock Exchange, 2024).	VenueBookState; NYSE ledger ranges
Hybrid priority	<code>rule737;</code> <code>rule737full</code>	Setter, displayed, and residual tiers compose and conserve quantity.	Setter eligibility, floor interest, complete rankings, and conditional instructions remain outside the kernel.	VenueBookState; NyseAuctionState
Displayed vs hidden	<code>t3_only_after_t2</code>	Reachable displayed interest is offered the residual before the modeled hidden tier.	Only two residual pools; not every venue-specific priority category.	VenueBookState
Protected quotes	<code>protected_;</code> <code>nbo</code>	The modeled NBO is no worse than any supplied protected offer; manual and sub-round-lot quotes do not set it.	Requires the correct venue-top set, security-specific round lot, accessibility, and data consolidation (U.S. Securities and Exchange Commission, 2005).	Rule 600 ledger entry; Rule611State
Rule 611	<code>tradeThroughBuy;</code> <code>sweepThenExecute</code>	No modeled trade-through at or below NBO; sweeps conserve and satisfy supplied better quotes.	Kernel omits timing, flicker, failures, and most exceptions (U.S. Securities and Exchange Commission, Division of Trading and Markets, 2008).	Rule611State; Rule611Exception
Self-help	<code>nboSH</code>	Removing a venue cannot improve the best remaining offer.	Systems-problem finding, notice, duration, policy, and resumption require evidence.	Rule 611 exception facts
Rule 610	<code>feeCap;</code> <code>accessFeeCap610;</code> <code>slideBuy</code>	Fee-spread bound; four dated cap branches; exact legacy dollar-price compatibility; sliding respects limit and avoids a modeled lock.	Fee schedules, rebate tiers, routing economics, and policies are not inferred from price arithmetic.	Dated FeeProfile610; September 2026 guard
Rule 612	<code>rule612;</code> <code>minimumTick612</code>	Legacy whole-cent kernel plus no-lock implies a one-cent spread; six dated tick branches; exact legacy predicate compatibility.	The amended profile remains delayed until the cited compliance date; measurement evidence is supplied.	Rule612State; schedule and September 2026 guard
Two clocks	<code>Ev</code>	Nonnegative latency, order preservation at equal latency, inversion possibility and bound.	Not SIP sequencing, synchronization, corrections, or feed arbitration (Ding et al., 2014 ; Bartlett and McCrary, 2019).	Analytic only; no claimed corpus analogue

Family	Kernel	Proved guarantee	Boundary	Clause-layer object
Order types	Midpoint, reserve, discretionary, peg, and post-only kernels	Floored midpoint is in [NBB, NBO); exact doubled midpoint is strictly inside; one-refresh conservation, peg monotonicity/no-lock, and non-taking properties.	The floored grid midpoint can equal the NBB at a one-unit spread; not a complete venue order life cycle or interaction model (The Nasdaq Stock Market LLC, 2026).	VenueBookState ; Nasdaq type/attribute catalogues
LULD	luldBands ; limitState_boundary	Static bands contain the reference and widen with the percentage.	Dynamic reference, tiers, timing, pauses, reopening, dissemination, and procedures are not kernel theorems.	LuldPlanState
Rule 201	shortSaleOK	While the supplied trigger is active, modeled short executions must be above NBB.	Trigger, duration, marking, policies, handling, and exceptions are not derived by this kernel.	Rule201State ; exception types
Rule 605	effSpread2	Signed $\mathbb{Z}$ kernel is bounded above by quoted spread, vanishes at an exact midpoint, and preserves below-midpoint magnitude.	Not by itself a side-aware reporting system; September and November profiles differ (U.S. Securities and Exchange Commission, 2024b ; U.S. Securities and Exchange Commission, Division of Trading and Markets, 2026 ; U.S. Securities and Exchange Commission, 2026c).	Rule605Report ; dated profile; exhaustive field types
Auctions	crossPrice3 ; dmmFacilitate	Lexicographic price key; sufficient facilitation clears residual and conserves.	Eligibility, imbalance publication, collars, freezes, reopening, floor procedures, and venue tie-breaks remain external.	NasdaqCrossState ; NyseAuctionState
State inference	Fingerprint, collision, and wheel-information theorems	Proves general nondegenerate phase separation, exact pointer necessity/sufficiency on common-observation fibers, reachable allowance recovery, and exact separation of mL complete-summary phases.	The mL bound omits side inputs; with aligned cumulative fills supplied, the reachable full-allowance fiber has exactly m pointer states and the allowance is redundant on every reachable live-head state.	Analytic only; reconstruction evidence is external
Conformance	conformsPT ; conformsPT_comp ; diagnose	Boolean equivalence to feasibility and seriality; first failure is named; individually conforming consecutive runs compose.	Price–time needs an authentic queue and order-to-fill alignment; parity additionally needs authenticated pointer state; neither is production-engine or transport certification.	Analytic only; production equivalence is external

14 Regulatory traceability: audit contract

This layer supports exactly three claims: the preceding source proves mathematical deductions, a compiler-linked ledger provides traceability into a fixed legal and venue corpus, and real-world compliance still requires independent legal, evidentiary, and implementation review. The proof standards are deliberately different. Appendix B contains the detailed profiles, source ranges, record types, and pinpoints.

The compiled corpus has 58 rows: 43 federal and federal-supporting ranges, five LULD Plan ranges, three NYSE ranges, and seven Nasdaq ranges. Each row has a unique citation, a readable label, and an elaborated Lean destination. Separate exhaustive-list checks cover the formal Rule 605 fields, Rule 611 exceptions, Rule 201 bases, Nasdaq order types, and Nasdaq attributes. These are audit lemmas about a finite encoding, not theorems that the encoding is a complete or correct legal interpretation.

The federal profiles use primary SEC materials available as of September 5, 2026. Release No. 34-105656 grants the stated relief for Rules 600(b)(89)(i)(F), 610(c), and 612 until the first business day of November 2027; Release No. 34-105655 remains expressly labelled a proposed rescission of Rules 611 and 610(e); Rule 605 FAQ 39 gives the November 1 collection date and end-of-December reporting date; its footnote 2 states ten order types per security; and Release No. 34-105136 is the April 1 Rule 605 exemptive order ([U.S. Securities and Exchange Commission, 2026b,a](#); [U.S. Securities and Exchange Commission, Division of Trading and Markets, 2026](#); [U.S. Securities and Exchange Commission, 2026c](#)).

Layer	Machine-checked result	Independent review still required
Numerical profile	Exact branch result for supplied inputs	Operative law, date, security class, normalization, and source facts
Finite catalogue or ledger	Every constructor or fixed row appears once with an elaborated destination	Completeness of the corpus and faithful interpretation of each source
Exception or attestation	Encoded conditions and required Boolean fields are present	Evidence, provenance, authenticity, legal availability, and implementation behavior

The ledger is therefore an inspection artifact attached to the mathematical paper, not evidence that any venue or broker-dealer complies with the mapped provisions.

15 Scope and limitations

Decided. The formal statements attached to the named Lean declarations: three integer allocation operators and their proved algebraic properties; the large-lot wheel-pass endpoint; price–time uniqueness; balance and an explicit run-dependent water level for the pass-based wheel; feasibility of an estate-indexed unit-increment rule and failure of its reset version to compose; one-lot balance and canonical stream composition for the same persistent pointer wheel, with no allowance premise on reachable states; impossibility of parity conformance from queue and aligned fills without pointer evidence; general nondegenerate pointer–allowance separation; exact fiberwise pointer necessity and sufficiency under a common authenticated observation; exact future equivalence, reachability, and the mL complete-summary lower bound for the finite phase family; the exact m pointer bound on the reachable full-allowance common-observation fiber; allowance recovery from cumulative allocation on every reachable live-head state; relabeling equivariance; signed effective-spread arithmetic; cross-market, order-type, auction, observability, and strategic kernels; and price–time conformance equivalence and sequential composition. Section 14 states the distinct fixed-corpus audit results. All 362 theorem declarations are kernel-checked, and the numerical worked examples are decided by evaluation. Fixed-length bit bounds are elementary deductions from the checked injectivity results. Corollary 6.5 is a conventional real-analysis consequence of the machine-checked finite approximation bound; it is not represented as a separate Lean declaration.

Not decided. Correctness of a legal interpretation beyond the stated formal mapping; truth, sufficiency, provenance, or authenticity of attestations and underlying evidence; behavioral equivalence between these specifications and a production exchange system; compliance by any real entity; provisions outside the stated fixed corpus; a global minimal quotient of every raw `WState` including claim-boundary and one-live-participant degeneracies; a characterization of all banded stream-consistent mechanisms; reconstruction of an order-level queue or participant alignment from consolidated data; empirical validity on exchange data; composition of the stateless reset estate rule (which is formally refuted); or continuous mixed-strategy equilibria. The executable price–time checker requires an authenticated or reconstructed order-level queue and queue-to-fill alignment and is exercised here only on synthetic inputs.

Security universe. The operator theorems do not depend on issuer capitalization or index membership. They therefore cover small-cap securities and Russell 3000 constituents whenever the stated queue, lot-size, venue, and state assumptions hold. The paper does not model index construction, constituent selection, capitalization, liquidity, or empirical performance.

16 Conclusion

Persistent wheel state restores split-versus-combined consistency lost under reset rotation. Future-fill probes establish pointer necessity on common-observation families, while the reachable phase invariant recovers unused allowance from aligned cumulative fills. These results connect executable allocation semantics to conformance and identifiability. Strategic consequences require additional behavioral assumptions; institutional application requires independent review of rule mappings, evidence, and implementation equivalence.

A The proof system

Lean 4.22.0, core library only, as a Lake package with no dependencies. The compiler identifies commit `ba2cbbf09d4978f416e0ebd1fceeabc2c4138c05` (Release). Modules in build order: `Rule737` (operators, large-lot endpoint, conservation, book share, worked allocation), `Markets` (three tiers, observability, cross, tick and protection, two clocks, routing), `Balance` (state-form wheel, explicit water level, balance theorem, two-level clearing key, fee cap), `Rulebook` (uniqueness, conformance, serial ladder, consolidation, sweep, sliding), `Orders` (order types, self-help, exact doubled midpoint, three-level clearing key, facilitation), `Strategic` (contests, dilution, routing), `SetterRace` (all-pay race), `Closure` (wheel agreement, ε -equilibrium), `Conformance` and `Collisions` (synthetic allocation records and collision witnesses), `Algebra` (operator monoid, conservation closure, impossibility, ladder uniqueness), `MoreRules` (completeness, estate completion/reset obstruction, signed execution-quality kernel, bands, close eligibility, post-only), `PathIndependence` (composition boundaries, canonical fuel, reachable phase invariant, idealized owner trace, pointer-relative balance), `WheelInformation` (general phase separation, fiberwise pointer minimality, finite-family reachability and exact quotients, reachable allowance recovery, relabeling equivariance, book-share counterexample), `ClauseComplete` (dated profiles, compatibility lemmas, 58-row ledger, finite schemas, exceptions, attestations, and checklist predicates), `Check` (exhaustive theorem-axiom audit), and `AllocationRecord` (executable checker). All 362 theorem declarations depend on at most `propext`, `Quot.sound`, and `Classical.choice`; worked examples depend on none.

Classical-choice audit. No source definition is marked `noncomputable`, and the source contains no call to `Classical.choose`. The executable `AllocationRecord` path uses `parseNats`, `firstDiff`, `diagnose`, `conformsPT`, and `priceTime`; none has an audited theorem dependency because definitions are evaluated rather than proved. Exactly 48 theorem proof terms include `Classical.choice`:

`wheelOp_conserve`, `wheelOpFull_conserve`, `rule737Op_conserve`, `ladder_unique`, `wheelS_within_L`, `wheelS_within_L_explicit`, `wheelS_eq_wheel`, `wheel_balanced`, `wheelS_state_eq`, `wheel_within_L`, `uniform_eps`, `uniform_eps_relative`, `priceTime_tail_invisible`, `priceTime_boundary_invisible`, `rule737full_sum`, `t3_only_after_t2`, `hidden_depth_lower_bound`, `supplyAt_zero_of_none`, `matched_dominated_by_ask_price`, `crossPrice_global`, `wheel_exhausts`, `discreteCEA_feasible`, `matchNYSE_conserve`, `nyse_price_priority`, `effSpread2_neg_iff`, `midpoint_inside`, `midpoint2_strict_inside`, `mpl_no_trade_through`, `zip_zero_add`, `priceTime_comp`, `conformsPT_comp`, `saturated_future_equiv_iff`, `saturated_future_sufficient_injective`, `runW_stops_of_tot_lt_fuel`, `runWFull_stopped`, `runWFull_allocates_first`, `runWFull_comp`, `runWFull_comp_reachable`, `wheel_sum`, `rule737_sum`, `nextWFill_crosses_allowance`, `validWheelPhase_probe_equiv_implies`, `wheelFiberSufficient_pointer_injective`, `phaseState_allowance_probe`, `phaseState_probe_equiv_iff`, `phaseState_future_equiv_iff`, `wheelPhaseSufficient_injective`, `phaseState_full_summary_injective`.

This proof dependency does not make the definitions non-executable.

A.1 Load-bearing theorem contracts

The table records the hypotheses that must travel with the headline results. It is a reading aid, not a substitute for the elaborated declarations.

Declaration	Required hypotheses / domain	Exact force of conclusion
<code>serial_unique</code>	Lists Q, A , quantity x ; <code>Feasible</code> $Q\ A\ x$; <code>Serial</code> $A\ Q$.	$A = \text{priceTime } Q\ x$. It does not authenticate Q or identify venue-specific queue construction.

Declaration	Required hypotheses / domain	Exact force of conclusion
<code>runW_band</code>	Natural L, n, x ; participant-labelled records R ; $0 < L$; state initialized by <code>initW</code> .	Some w makes the returned participant state satisfy Band L w . It is fuel-bounded and says nothing about eligibility inputs.
<code>runW_comp</code>	First run has zero leftover; second result satisfies Stopped ; combined run receives fuel $n_1 + n_2$.	Split and combined runs return exactly the same rotated state and leftover. No claim without the completion and termination premises.
<code>runWFull_comp</code>	$0 < L$; positive current allowance; first estate $x \leq \text{totW}(st.S)$; canonical fuel $\text{totW}(st.S) + 1$.	For every y , combined and split canonical runs are exactly equal. Completion, stopping, and fuel are derived rather than supplied.
<code>runWFull_comp_reachable</code>	State returned by <code>runW</code> from <code>initW</code> ; $0 < L$; first estate $x \leq \text{totW}(st.S)$.	The same exact split/combined equality, with current-allowance positivity derived from reachability.
<code>future_sufficient_summary_requires_pointer_information</code>	Arbitrary summary codomain; a predictor correct for every four-party state and next quantity from canonical queue, cumulative fills, summary, and quantity.	The summary must distinguish pointerA from pointerC . This is a witness-level information lower bound, not global minimality of WState .
<code>validWheelPhase_probe_equiv_implies</code>	Two states with distinct live head/successor labels, allowances in $[1, L]$, and every remaining claim above L .	Equality of all participant-labelled probes through L forces equal pointer labels and allowances; other state components are not thereby identified.
<code>wheelFiberSufficient_pointer_injective</code> ; <code>wheelPointerCode_fiber_sufficient</code>	Any one-state-per-pointer nondegenerate family; $L > 0$; an authenticated observation common across that family; one decoder correct for every indexed state, estate, and label.	Every future-sufficient summary is injective in the pointer, and the head label is sufficient for the fixed family. This is exact fiberwise minimality, not a global quotient of arbitrary raw states.
<code>phaseState_future_equiv_iff</code> ; <code>wheelPhaseSufficient_injective</code>	Reachable family $P_{p,a}$; $m \geq 2$; $p < m$; $1 \leq a \leq L$; complete summary with no separate queue/fill input.	Actual <code>runWFull</code> probes of size at most L separate every phase. The pair is sufficient. Counting yields mL classes and the binary bit bound.
<code>phaseState_full_summary_injective</code> ; <code>phaseState_full_pointer_sufficient</code>	Reachable fresh-lot states $P_{p,L}$; $m \geq 2$; $L > 0$; aligned cumulative-allocation observation supplied.	The common zero observation still requires m distinct summary values, and the pointer supplies exactly m . The bit bound is the elementary encoding corollary.
<code>runW_allowance_phase</code>	State returned by <code>runW</code> from <code>initW</code> ; $L > 0$; returned state has a live head with cumulative allocation A_p .	Its allowance is $L - (A_p \bmod L)$. The cycle-phase equation and allowance positivity are derived rather than supplied.
<code>runWFull_relabel</code>	Relabel every identifier consistently; preserve the ordered records and allowance.	The canonical returned state is relabeled and the leftover unchanged. For bijective labels this preserves a valid participant set.
<code>checkedRoundLotSize600_off_grid</code>	Existing stock; price input in \$0.0001 units; input not divisible by 100.	Returns no tier. Normalization and the legal evaluation/operative period require independent evidence.
<code>discreteCEA_feasible</code> ; <code>discreteCEA_reset_not_composable</code>	Any natural claims and estate for feasibility; fixed claims $(10, 10)$ for the counterexample.	The unit rule completes every integer estate, but deterministic rotation reset violates composition: two one-unit estates differ from one two-unit estate.
<code>waterFill_two_equal_no_single_share</code>	Any natural water level s ; fixed claims $(10, 10)$.	Benchmark award sum is not 1. Therefore natural water levels do not parameterize every integer estate.

Declaration	Required hypotheses / domain	Exact force of conclusion
<code>conformsPT_iff</code>	Supplied queue Q , aligned fill vector A , and quantity x .	Boolean acceptance iff A is feasible and serial relative to those supplied inputs. No authenticity or public-tape reconstruction guarantee.
<code>midpoint2_strict_inside</code>	Natural $nbb < nbo$.	$2nbb < nbb + nbo < 2nbo$ in doubled units. The integer-grid <code>midpoint</code> theorem is only half-open.
<code>clause_corpus_has_58_entries</code>	The fixed compiled <code>clauseCorpus</code> .	Its list length is 58; this is an audit lemma, not semantic or legal completeness.

A.2 Paper and source availability

The paper’s DOI is [10.5281/zenodo.22343804](https://doi.org/10.5281/zenodo.22343804).

The Lean source code, proof scripts, and reproducible build instructions will be made publicly available shortly at <https://github.com/MiquelNoguerAlonso/Formal-Verification>. The release will include the LaTeX sources, a deterministic build script, and a SHA-256 manifest for the distributed files. No empirical dataset is used.

B Dated regulatory traceability and inspection

This section supports exactly three distinct claims: the preceding source proves mathematical deductions, the clause ledger provides compiler-linked traceability into a fixed corpus, and real-world compliance still requires independent legal, evidentiary, and implementation review. The allocation operators and rule-inspection records share one formal development, but the evidence standard does not. The profiles and mappings below make the second claim reviewable without promoting it into either of the other two.

The traceability unit is an expressly cited paragraph, paragraph range, or identified Commission guidance/order item, not necessarily one row per printed subparagraph. A compiler-linked destination ensures that a referenced formal object exists and survives refactoring. It does not establish that the destination faithfully or sufficiently interprets the source.

B.1 Fixed and dated corpus

The 58 rows comprise 43 federal and federal-supporting ranges, five LULD Plan ranges, three NYSE ranges, and seven Nasdaq ranges. The federal snapshot is the eCFR text current through September 2, 2026, read with operative-date, guidance, and exemptive-relief materials; venue texts were accessed September 4, 2026 ([Office of the Federal Register and U.S. Government Publishing Office, 2026a,b](#); [New York Stock Exchange, 2026](#); [The Nasdaq Stock Market LLC, 2026](#); [Limit Up-Limit Down Plan Participants, 2026](#)). The corpus is:

- three Rule 600(b) rows—general incorporated definitions, paragraph (89)(i)(F) minimum-increment indicator with its June 2026 relief profile, and paragraph (93) round-lot tiers—plus eight Rule 610 ranges, six Rule 611 ranges, seven Rule 612 ranges, nine Rule 201 ranges, and ten Rule 605 rule/guidance/order ranges, so $3 + 8 + 6 + 7 + 9 + 10 = 43$;
- five operative LULD Plan ranges grouping definitions and coverage, reference and numerical rules, band and pause states, reopening and halts, and dissemination/policy/overnight duties;
- NYSE Rules 7.31(a)–(i), 7.35 general/A/B, and 7.37(a)–(g); and
- Nasdaq Rules 4702(a), 4702(b)(1)–(16), 4703(a)–(l), 4752(a)–(f), 4753(a)–(f), 4754(a)–(c), and 4757(a)–(c).

These counts describe mapping granularity. They are not a count of every nested sentence, proviso, cross-reference, or provision elsewhere in the Exchange Act, Regulation NMS, Regulation SHO, or the complete venue rulebooks. Completeness is claimed only for finite formal catalogues whose constructors and advertised lists are explicitly compared by Lean.

B.2 Formalization of non-numerical duties

A numerical clause is executable. For example, `accessFeeCap610` selects the fee cap from an explicit profile and price; `minimumTick612` selects the increment from profile, price, spread, and new-stock status; and `tenPercentDecline201` evaluates the short-sale trigger.

A legal exception is represented by a finite reason together with the facts that validate it. `Rule611Exception` has one constructor for each Rule 611(b) exception. `ShortExemptReason201` and `ShortExemptFacts201` encode the Rule 201 short-exempt bases and their internal conditions. Whole-record relief is not a bypass switch: `ExemptionAttestation` separately records a grant,

coverage of every checked duty, and satisfaction of stated conditions. Partial relief must be represented at the affected duty.

Procedural language such as establishing written policies becomes a named Boolean attestation. That makes the premise visible and machine-checkable, but the Boolean is not evidence. It does not authenticate manuals, notices, surveillance records, market data, or production behavior. Provenance and independent verification remain external.

B.3 Federal specifications

B.3.1 Rule 600 definitions

General incorporated definitions exclude paragraphs (89)(i)(F) and (93), which have dedicated rows. Paragraph (89)(i)(F) requires an indicator of the applicable minimum pricing increment under Rule 612; it is not the odd-lot-information definition. Its `MinimumIncrementIndicatorProfile600` selects the June 11, 2026 relief regime for the September snapshot. Odd-lot information is defined in paragraph (69), including the best odd-lot additions in (69)(iii), and has a distinct dissemination timetable. The adopted text makes this distinction on printed pages 512–513 of Release No. 34-101070; the indicator’s relief is stated on pages 5–6 and 8 of Release No. 34-105656 ([U.S. Securities and Exchange Commission, 2024a, 2026b](#)).

Paragraph (93) has the numerical `roundLotSize600` kernel: 100 shares through \$250, 40 through \$1,000, 10 through \$10,000, one above that, and 100 for a new NMS stock. The inspected interface `checkedRoundLotSize600` accepts an existing-stock average supplied in normalized cents, represented in the paper’s \$0.0001 units, and rejects off-grid inputs. It therefore does not silently classify \$250.0001: that input returns `none`, whereas \$250.01 returns 40. Five tier lemmas and four interface/boundary lemmas verify the numerical behavior. The input’s normalization, primary-exchange provenance, and assignment to the March/September evaluation and May/November operative periods remain explicit external duties. The interface does not prescribe a rounding policy for raw averages; paragraph (93) and its operative timetable are on printed pages 513–514 of the adopting release ([U.S. Securities and Exchange Commission, 2024a](#)).

B.3.2 Rule 610

`Rule610State` separately records fair access to an SRO trading facility, equivalent and nondiscriminatory access to a display-only facility, the price-dependent fee cap, fee transparency at execution, lock/cross controls, pattern-or-practice control, and scoped Commission relief. `FeeProfile610` distinguishes the legacy cap used during the September 2026 relief period from the lower amended cap. Four branch theorems prove the dollar and sub-dollar formulas. `rule610Complies` rejects the amended profile for the dated snapshot; `rule610_amended_profile_not_september2026` proves that guard ([U.S. Securities and Exchange Commission, 2024a, 2026b](#)). The Commission’s June 11, 2026 action concerning Rules 611 and 610(e) was a rescission proposal, not a final rescission; the fixed corpus therefore retains both operative specifications and labels the proposal as such ([U.S. Securities and Exchange Commission, 2026a](#)).

B.3.3 Rule 611

`Rule611State` requires written policies, regular surveillance, prompt remediation, transaction-level protection, ISO reasonable steps, and any scoped relief. `Rule611Exception` exhausts system failure, non-regular-way settlement, single-price auctions, crossed protected markets, incoming ISO, outbound ISO sweep, benchmark trades, the one-second quotation condition, and stopped

orders. The stopped-order case additionally requires customer status, order-by-order agreement, and the prescribed price relation. `rule611Complies` rejects a trade-through unless a fully conditioned alternative holds.

B.3.4 Rule 612

`Rule612State` records the pricing profile, evaluation and operative periods, primary-listing-exchange measurement, covered price, new-stock status, display, ranking, acceptance, indications of interest, and scoped relief. `operativePeriodFor612` maps January–March measurements to May–October and July–September measurements to November–April. `minimumTick612` encodes both the legacy-relief profile and the delayed amended profile. Six theorems prove the profile/price branches; `rule612_iff_minimumTick612_legacy` proves agreement with the analytic legacy predicate. The September guard prevents the delayed amended profile from being reported as the current decision (U.S. Securities and Exchange Commission, 2026b).

B.3.5 Rule 201

`restrictionApplies201` derives the ten-percent decline condition, requires listing-market determination and notice, restricts the remainder of that day and the following day, and requires a continuously disseminated NBB. `transactionAllowed201` implements the above-NBB, initially displayed, and short-exempt alternatives. `validShortExempt201` checks ownership/delivery restriction, odd-lot market making, convertible and foreign arbitrage, underwriting cases, riskless-principal sequencing/allocation/records, and VWAP conditions. Policies, surveillance, remediation, SRO conformity, and valid scoped relief complete `rule201Complies`.

B.3.6 Rule 605

The amended schema is finite and encoded without ellipses. `OrderType605` has ten covered order classes. `CoreField605`, `ImmediateField605`, `NonmarketableField605`, and `SummaryField605` contain 25, 19, 4, and 12 fields. Five `*605_complete` theorems prove by exhaustive case analysis that each full-schema list contains every constructor.

The general amended rule began August 1, 2026. SEC Rule 605 FAQ 39, answer A.39, states that collection of the five best-available-displayed-price statistics begins November 1, 2026 and that the first reports containing them are due at the end of December 2026; FAQ footnote 2 identifies ten order types per security (U.S. Securities and Exchange Commission, Division of Trading and Markets, 2026). `Rule605Profile` therefore distinguishes `augustThroughOctober2026` from `november2026AndLater`. The first profile requires 14 immediate fields and the second all 19. Theorems prove both branches and reject the 19-field profile for September. `Rule605Report` also records the treatments in Sections II.A–II.B of the April 1, 2026 order, Exchange Act Release No. 34-105136 (U.S. Securities and Exchange Commission, 2026c). `publicationRoute605` uses uniform plan access when an effective joint plan exists and the free, machine-readable website fallback otherwise.

B.4 Plan and venue specifications

B.4.1 Limit Up–Limit Down

`LuldPlanState` names stock coverage; tier and price classification; rolling reference and change threshold; percentages, late-day doubling, and rounding; execution and quotation bands; limit and

straddle states; the pause; reopening and regulatory halts; processor dissemination; participant policies and surveillance; overnight protection; and exemption. `luldPlanComplies` permits no silent omission.

B.4.2 Nasdaq

`NasdaqOrderType4702` has all sixteen Rule 4702 order types and `NasdaqAttribute4703` all twelve Rule 4703 attribute families. Their exhaustiveness theorems establish list completeness. `VenueBookState` covers session eligibility, type semantics, attribute compatibility, Reg NMS and Reg SHO repricing, timestamp behavior, priority, routing, data and fallback sources, locks, self-help, cancellation, and re-entry. `NasdaqCrossState` covers definitions, eligible interest, imbalance data and schedules, price hierarchy, collars, allocation priority, uniform price, official-price dissemination, LULD/hybrid variants, and contingencies.

B.4.3 NYSE

The shared book-state record covers Rule 7.31 order/modifier behavior and Rule 7.37 execution, allocation, routing, market data, lock/cross, and self-help duties; `rule737full` and `matchNYSE` are the proved allocation kernels in the foundations paper. `NyseAuctionState` represents Rule 7.35 general, Core Open/Trading Halt, and DMM-facilitated Closing Auction families, including responsibilities, eligibility, entry/modification/cancellation windows, imbalance publication, reference prices, collars/freezes, price determination, priority, facilitation, reopening, and contingencies.

B.5 Ledger audit and exact force

The 58-entry `clauseCorpus` maps each identified source unit to a `FormalDestination` and classifies it as a definition, numerical rule, state transition, exception, procedure, or disclosure. A destination stores an actual Lean value, function, projection, or tuple; deleting or renaming the referenced object breaks elaboration.

Audit lemma B.1 (Fixed-corpus checks). *The declarations `clause_corpus_has_58_entries`, `clause_corpus_family_counts`, `clause_corpus_citations_unique`, and `clause_corpus_traceability_complete` verify, respectively, the total count, the advertised $43 + 5 + 3 + 7$ partition, unique citations, and that every row has readable labels and a compiler-linked destination. They are list-length and elaboration audits, not substantive legal theorems.*

Generated ranges are not deemed complete merely because they have a destination. Rule 605 fields, Rule 611 exceptions, Rule 201 bases, Nasdaq order types, and Nasdaq attributes are finite inductive types compared with exhaustive lists. These audit lemmas do not prove faithful interpretation, sufficient grouping, evidentiary truth, implementation equivalence, or venue compliance.

B.6 Inspection contract

Layer	Machine-checked result	Independent review still required
Numerical profile	Exact branch result for supplied date/profile inputs	Correct operative law, date, security class, and source facts
Finite catalogue	Every formal constructor appears in the advertised list	Whether the constructors faithfully partition the controlling text

Layer	Machine-checked result	Independent review still required
Exception	Enumerated reason satisfies encoded internal conditions	Evidence for those facts and legal availability of the exception
Attestation	Required Boolean fields and scoped-relief conditions are present	Provenance, authenticity, sufficiency, and reviewer judgment
Clause ledger	58 unique rows have elaborated destinations and labels	Completeness outside the fixed corpus and semantic quality of each mapping

C Compiler-checkable source excerpt

The definitions below are an ASCII transcription of the corresponding Lean source; `Prod` denotes an ordered pair. The final `#check` commands resolve thirteen declarations, including the general phase-separation, fiberwise pointer-minimality, and reachable allowance theorems.

```

import Rulebook
import Balance
import PathIndependence
import WheelInformation
import Strategic
namespace PaperExcerpt

def give (L r x : Nat) : Nat := min L (min r x)

def wheelRound (L : Nat) : List Nat -> Nat -> Prod (List Nat) Nat
| [], x => ([], x)
| r :: R, x =>
  let g := give L r x
  let p := wheelRound L R (x - g)
  (g :: p.1, p.2)

def wheel (L : Nat) : Nat -> List Nat -> Nat -> Prod (List Nat) Nat
| 0, R, x => (R.map (fun _ => 0), x)
| n + 1, R, x =>
  if x = 0 then (R.map (fun _ => 0), x)
  else
    let p := wheelRound L R x
    let R' := List.zipWith (fun r g => r - g) R p.1
    let q := wheel L n R' p.2
    (List.zipWith (fun a b => a + b) p.1 q.1, q.2)

def priceTime : List Nat -> Nat -> List Nat
| [], _ => []
| q :: Q, x => min q x :: priceTime Q (x - min q x)

#check @Rule737.serial_unique
#check @Rule737.wheelS_band
#check @Rule737.runW_band
#check @Rule737.runWFull_comp_reachable
#check @Rule737.tullock_best_response
#check @Rule737.validWheelPhase_probe_equiv_implies
#check @Rule737.wheelFiberSufficient_pointer_injective
#check @Rule737.phaseState_future_equiv_iff
#check @Rule737.phaseState_reachable
#check @Rule737.wheelPhaseCode_sufficient
#check @Rule737.phaseState_full_summary_injective
#check @Rule737.runW_allowance_phase
#check @Rule737.runWFull_relabel
end PaperExcerpt

```

D Certified inequalities for the strategic layer

The following payoff comparisons use the stated allocation rules and additional behavioral assumptions. They are expressed with denominators cleared over $\mathbb{Z}$.

D.1 Contests for a slot or a tranche

m contestants with linear effort cost compete for a prize V ; effort e' against others' total O earns $\pi(e') = Ve'/(e' + O) - e'$ (Tullock, 1980). Write $D = e + O$, $t = e' + O$.

Theorem D.1 (Best response; `tullock_best_response`). *If $D = e + O$, $VO = D^2$, $t = e' + O > 0$ and $D > 0$, then $(Ve' - e't)D \leq (Ve - eD)t$, i.e. $\pi(e') \leq \pi(e)$ for every $e' \geq 0$.*

Proof idea. $Dt[\pi(e) - \pi(e')] = D(D - t)^2$ under $VO = D^2$: a positive number times a square. $\square$

Theorem D.2 (Symmetric equilibrium and dissipation; `tullock_symmetric_foc`, `tullock_dissipation`, `dissipation_increasing`). *If $m^2e = (m - 1)V$ then the first-order condition holds at $O = (m - 1)e$, total effort satisfies $m \cdot me = (m - 1)V$, and $(m - 1)(m + 1) < m^2$: a fraction $(m - 1)/m$ of the prize is dissipated and is strictly increasing in m .*

Example D.3 (`ex_tullock`, `ex_tullock_dev`). $m = 4$, $V = 1200$: $e = 225$, $D = 900$, $O = 675$, $VO = 810,000 = D^2$; the deviation $e' = 300$ is unprofitable.

D.2 The race for setter priority

Two racers invest speed $b_1, b_2 \geq 0$ (sunk) for a tranche S ; the faster wins, ties split. Doubled payoffs keep the model in $\mathbb{Z}$ (`u2`).

Theorem D.4 (`setter_race_no_pure_NE`, `race_spend`). *For $S \geq 3$ no profile is a pure equilibrium: ties invite a one-unit raise, gaps of two or more invite undercutting, and a gap of one invites the loser to drop out or leapfrog. In every profile the doubled payoffs sum to $2S - 2(b_1 + b_2)$.*

Theorem D.5 (Uniform mixture is an ε -equilibrium; `sumU`, `sumU_closed`, `uniform_payoff`, `uniform_eps`, `uniform_selfpayoff_zero`, `uniform_eps_relative`). *On the grid $\{0, \dots, N\}$ with prize N and a uniform opponent, the $(N + 1)$ -scaled doubled payoff of bid b is $N - 2b$; the uniform strategy's own scaled payoff is 0; the best pure deviation improves expected payoff by less than half a cost unit, and relative to the prize the gain is less than $1/N$.*

These identities give the finite-grid dissipation result for the all-pay auction (Baye et al., 1996).

D.3 Privileged slot and venue competition

Theorem D.6 (Privileged-slot comparative statics). Lean: `dmm_share_antitone`, `dmm_option_dilution`, `dmm_option_spread`, and `dmm_advantage`. *A free parity slot worth $\lfloor u/m \rfloor$ units times the half-spread h is non-increasing in m and non-decreasing in h ; relative to a participant paying c per slot, the advantage is exactly c .*

Theorem D.7 (Routing; `bestVenue_le`, `flowShare`, `undercut_wins`, `fee_cut_captures`, `fee_gap_bounded`). *The router's chosen net cost is at most every venue's; the venue with strictly lower net cost takes all marketable flow; a fee cut below the rival's net cost captures it; under the fee cap the net-cost gap at equal quoted prices is at most the cap.*

Principle D.8 (Layering). The strategic layer introduces additional behavioral assumptions. Its propositions are proved against the stated operators, while every translation from institutional rules to those operators remains a separately documented modelling assumption.

References

- Robert J. Aumann and Michael Maschler. Game theoretic analysis of a bankruptcy problem from the talmud. *Journal of Economic Theory*, 36(2):195–213, 1985.
- Patrick Bahr, Jost Berthold, and Martin Elsman. Certified symbolic management of financial multi-party contracts. In *ICFP 2015*, pages 315–327. ACM, 2015.
- Robert P. Bartlett and Justin McCrary. How rigged are stock markets? evidence from microsecond timestamps. *Journal of Financial Markets*, 45:37–60, 2019.
- Michael R. Baye, Dan Kovenock, and Casper G. de Vries. The all-pay auction with complete information. *Economic Theory*, 8(2):291–305, 1996.
- Paul Brennan. An introduction to Imandra Markets. Imandra Markets technical case study; <https://medium.com/imandra/an-introduction-to-imandra-markets-d9b7419916f2>, 2024. Published February 20, 2024; describes exchange digital twins, matching-logic changes, design verification, and model-based validation.
- Eric Budish, Peter Cramton, and John Shim. The high-frequency trading arms race: Frequent batch auctions as a market design response. *Quarterly Journal of Economics*, 130(4):1547–1621, 2015.
- Carole Comerton-Forde, Vincent Grégoire, and Zhuo Zhong. Inverted fee structures, tick size, and market quality. *Journal of Financial Economics*, 134(1):141–164, 2019.
- Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *Automated Deduction – CADE 28*, volume 12699 of *LNCS*, pages 625–635. Springer, 2021.
- Shengwei Ding, John Hanna, and Terrence Hendershott. How slow is the NBBO? a comparison with direct exchange feeds. *Financial Review*, 49(2):313–332, 2014.
- Mohit Garg and Suneel Sarswat. The design and regulation of exchanges: A formal approach. In *42nd IARCS Annual Conference on Foundations of Software Technology and Theoretical Computer Science*, volume 250 of *Leibniz International Proceedings in Informatics*, pages 39:1–39:21. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik, 2022. doi: 10.4230/LIPIcs.FSTTCS.2022.39.
- Lawrence Harris. *Trading and Exchanges: Market Microstructure for Practitioners*. Oxford University Press, 2003.
- Carmen Herrero and Ricardo Martínez. Balanced allocation methods for claims problems with indivisibilities. *Social Choice and Welfare*, 30(4):603–617, 2008. doi: 10.1007/s00355-007-0262-z.
- Carmen Herrero and Antonio Villar. The three musketeers: four classical solutions to bankruptcy problems. *Mathematical Social Sciences*, 42(3):307–328, 2001.
- Limit Up-Limit Down Plan Participants. Limit up-limit down plan and amendments. <https://www.luldplan.com/plans>, 2026. Plan materials accessed September 4, 2026.
- Denis Merigoux, Nicolas Chataing, and Jonathan Protzenko. Catala: A programming language for the law. In *ICFP 2021, Proc. ACM Program. Lang.* 5, 2021.
- Hervé Moulin. Priority rules and other asymmetric rationing methods. *Econometrica*, 68(3):643–684, 2000.
- John Nagle. On packet switches with infinite storage. *IEEE Transactions on Communications*, 35(4):435–438, 1987.
- Raja Natarajan, Suneel Sarswat, and Abhishek Kr Singh. Verified double sided auctions for financial markets. In *12th International Conference on Interactive Theorem Proving*, volume 193 of *Leibniz International Proceedings in Informatics*, pages 28:1–28:18. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik, 2021. doi: 10.4230/LIPIcs.ITP.2021.28.
- Anil Nerode. Linear automaton transformations. *Proceedings of the American Mathematical Society*, 9(4):541–544, 1958. doi: 10.1090/S0002-9939-1958-0135681-9.
- New York Stock Exchange. NYSE parity: Frequently asked questions. https://www.nyse.com/publicdocs/nyse/NYSE_Parity_FAQs.pdf, 2024. Accessed September 4, 2026.
- New York Stock Exchange. NYSE rules: All NYSE group exchanges. <https://www.nyse.com/regulation/rules>, 2026. Current-rule access point, accessed September 4, 2026.
- Office of the Federal Register and U.S. Government Publishing Office. Electronic code of federal regulations, title 17, part 242: Regulation NMS. <https://www.ecfr.gov/current/title-17/chapter-II/part-242/subject-group-ECFRac68bdd026a46db>, 2026a. Up to date through September 2, 2026; accessed September 4, 2026.
- Office of the Federal Register and U.S. Government Publishing Office. Electronic code of federal regulations, title 17, part 242, section 242.201: Circuit breaker. <https://www.ecfr.gov/current/title-17/chapter-II/part-242/subject-group-ECFR1607681c7b4f78d/section-242.201>, 2026b. Up to date through September 2, 2026; accessed September 4, 2026.
- Maureen O’Hara. *Market Microstructure Theory*. Blackwell, 1995.

- Maureen O'Hara and Mao Ye. Is market fragmentation harming market quality? *Journal of Financial Economics*, 100(3):459–474, 2011.
- Christine A. Parlour. Price dynamics in limit order markets. *Review of Financial Studies*, 11(4):789–816, 1998.
- Grant O. Passmore and Denis Ignatovich. Formal verification of financial algorithms. In *Automated Deduction – CADE 26*, volume 10395 of *LNCS*, pages 26–41. Springer, 2017.
- Suneel Sarswat and Abhishek Kr Singh. Formally verified trades in financial markets. In *Formal Methods and Software Engineering*, volume 12531 of *Lecture Notes in Computer Science*, pages 217–232. Springer, 2020. doi: 10.1007/978-3-030-63406-3_13.
- The mathlib Community. The Lean mathematical library. In *Proceedings of CPP 2020*, pages 367–381. ACM, 2020.
- The Nasdaq Stock Market LLC. Nasdaq Equity 4: Equity rules. <https://listingcenter.nasdaq.com/rulebook/nasdaq/rules/Nasdaq%20Equity%204>, 2026. Rulebook snapshot accessed September 4, 2026.
- William Thomson. Axiomatic and game-theoretic analysis of bankruptcy and taxation problems: a survey. *Mathematical Social Sciences*, 45(3):249–297, 2003.
- Gordon Tullock. Efficient rent seeking. *Toward a Theory of the Rent-Seeking Society* (Buchanan, Tollison, Tullock, eds.), *Texas A&M University Press*, pages 97–112, 1980.
- U.S. Securities and Exchange Commission. Regulation NMS. Technical Report Release No. 34-51808, SEC, 2005. URL <https://www.sec.gov/files/rules/final/34-51808fr.pdf>.
- U.S. Securities and Exchange Commission. Regulation NMS: Minimum pricing increments, access fees, and transparency of better priced orders. Release No. 34-101070; <https://www.sec.gov/files/rules/final/2024/34-101070.pdf>, 2024a.
- U.S. Securities and Exchange Commission. Disclosure of order execution information. Release No. 34-99679; <https://www.sec.gov/rules-regulations/2024/03/disclosure-order-execution-information>, 2024b. Final rule issued March 6, 2024.
- U.S. Securities and Exchange Commission. Order granting temporary exemptive relief from compliance with certain rules under regulation NMS. <https://www.sec.gov/newsroom/press-releases/2025-130-sec-issues-exemptive-order-regarding-compliance-certain-rules-under-regulation-nms>, 2025. Issued October 31, 2025.
- U.S. Securities and Exchange Commission. Proposed rescission of regulation NMS rules 611 and 610(e). Release No. 34-105655, File No. S7-2026-20; <https://www.sec.gov/rules-regulations/2026/06/s7-2026-20>, 2026a. Official SEC page labels the action “Proposed”; issued June 11, 2026, published June 17, and comments due August 17, 2026; no final rescission identified as of September 5, 2026.
- U.S. Securities and Exchange Commission. Order granting temporary exemptive relief from compliance with rules 600(b)(89)(i)(f), 610(c), and 612 of regulation NMS, as amended. Release No. 34-105656; <https://www.sec.gov/files/rules/exorders/2026/34-105656.pdf>, 2026b. Issued June 11, 2026; pp. 5–6 and 8: relief for the minimum-pricing-increment indicator in 600(b)(89)(i)(F), Rule 612, and amended access-fee caps until the first business day of November 2027; odd-lot information is separately identified as 600(b)(69) on p. 2.
- U.S. Securities and Exchange Commission. Order granting exemptive relief from rule 605 of regulation NMS. Release No. 34-105136; <https://www.sec.gov/files/rules/exorders/2026/34-105136.pdf>, 2026c. Issued April 1, 2026; Sections II.A–II.B; p. 1 traces the general compliance date to August 1, 2026; verified September 5, 2026.
- U.S. Securities and Exchange Commission, Division of Trading and Markets. Responses to frequently asked questions concerning rules 610, 611, and 612 of regulation NMS. <https://www.sec.gov/divisions/marketreg/nmsfaq610-11.htm>, 2008. Staff guidance, accessed September 4, 2026.
- U.S. Securities and Exchange Commission, Division of Trading and Markets. Frequently asked questions: Rule 605 of regulation NMS. <https://www.sec.gov/rules-regulations/staff-guidance/trading-markets-frequently-asked-questions/frequently-asked-questions-rule-605-regulation-nms>, 2026. Staff guidance dated April 1, 2026; FAQ 39, answer A.39, states November 1, 2026 for collection and end-December 2026 for first reports; footnote 2 states ten order types per security; verified September 5, 2026.
- Chen Yao and Mao Ye. Why trading speed matters: A tale of queue rationing under price controls. *Review of Financial Studies*, 31(6):2157–2183, 2018.
- H. Peyton Young. On dividing an amount according to individual claims and liabilities. *Mathematics of Operations Research*, 12(3):398–414, 1987.
- H. Peyton Young. *Equity: In Theory and Practice*. Princeton University Press, 1994.